

# Building Robust Text Processing Pipelines

Maryam Jahanshahi Ph.D.

# Text Analysis: The Final Frontier

Extracting structured insights from unstructured inputs is not easy

The amount of  
digital data is  
**doubling**  
every year

Less than  
**0.5%**  
of all data in a  
company is  
analyzed

More than  
**75%**  
of all data  
is in text form



# About me

Past

**Journalist > Editor**  
on pipeline building

**Cancer Biologist**  
on scaling sizes

Present

**Research Scientist**  
TapRecruit

How do we make decisions  
about our careers?

What role do documents play  
in this process?

Future



# Achtung!

The focus of this talk will **not** be about text analysis in R

## Corpus Processing

**spaCy**

## Language Modeling

**gensim**

**TensorFlow**

## Stack Considerations

- Actively developed libraries
- Industrial-strength NLP:
  - Parallel processing of large datasets
  - Prototype to Production
  - Effective memory management

## Design Considerations

- Maintainability
- Reproducibility
- Environments

# A Day in the Life of an NLP Project



# A Day in the Life of an NLP Project



**Corpus:**  
A collection  
of documents



**Document:**  
Unprocessed  
string, typically  
associated with  
structured data

'Amanda  
didn't  
start the fire'

**Segment:**  
Processed  
string  
(i.e. sentence,  
paragraph etc.)

('fire', NN)

**Token:**  
Processed  
single data  
point



# Designing Data Preprocessors

Clean up

**Goal:** Remove inconsistency between otherwise similar data points

Segmentation

**Goal:** Split text chunks into data points (i.e. the unit of analysis or evaluation)

Normalization

**Goal:** Put data points on an equal footing

# Designing Data Preprocessors

Clean up

Segmentation

Normalization

## General Considerations

- What is the unit or data structure of analysis?  
(Tokens vs sentences vs paragraphs vs docs)
- Can the cleanup aid segmentation?

## Specific Considerations

- What is the role of punctuation?
- What role do hyphenated words play?
- Will parsing emojis or emoticons be helpful?



# Functions in a typical clean up script

The order of operations is important

HTML/XML

<p>Amanda didn't start the fire!</p>

Python: **Beautiful Soup**  
R: **xml2?**

Unicode

<p>Amanda didn't start the fire!</p>

Python: **regex**  
R: **utf8**

Contractions

Amanda did not start the fire!


Python: **spaCy**  
R: **textclean**

Punctuation

Amanda did not start the fire

Python: **spaCy**  
R: **textclean**

# Word incidence is rarely distributed normally



- **Stop words:** Removing most frequent words.
  - Standard list with most NLP libraries
  - Make your own artisanal list
- **Changing cases:**
  - Standard is to convert to lower case
  - Casing may matter for you (e.g. IT vs it)
- **Process numbers:**
  - Standard is to remove all numbers
  - Convert into words via **inflect** library in Python and **textclean** package in R
- **Stem or lemmatize words:**
  - Standard is to lemmatize

# Best Practices in Data Organization

```
corpus
|_README.md
|_raw
|   |_01.txt
|   |_02.txt
|   |_03.txt
|   |_metadata.json
|_processed
|   |_processed.json
|   |_metadata.json
|_scripts
```

## **Processed documents:**

Save down processed documents either as JSON objects or in a document database (NoSQL)

## **Metadata:**

Define what has been processed and when in metadata:

- Files
- Words
- Unique Tokens
- Date of latest preprocessing

# Advanced Best Practices in Data Organization

## Create a corpus reading module:

- Define which files should be loaded and how those files should be loaded.
  - Store these as parameters in README.
  - Regex for file names / formats [**\w\.****txt+**]
  - Can include a filter list for restricting files

```
corpus
|_README.md
|_raw
|   |_01.txt
|   |_02.txt
|   |_03.txt
|   |_metadata.json
|_processed
|   |_processed.json
|   |_metadata.json
|_scripts
```



# Advanced Best Practices in Data Organization

## Create a corpus reading module:

- Define which files should be loaded and how those files should be loaded.
  - Store these as parameters in README.
  - Regex for file names / formats [**\w\.**txt+]
  - Can include a filter list for restricting files

```
import json

def project_reader(self):
    return json.load(self.open("README"))
```

# Feature Extraction

Bag of Words representation vectorizes through word counts

	amanda	baby	did	doo	fire	not	shark	start	the
--	--------	------	-----	-----	------	-----	-------	-------	-----

Amanda didn't start the fire

1	0	1	0	1	1	0	1	1
---	---	---	---	---	---	---	---	---

Baby shark, doo doo doo doo doo doo

0	1	0	6	0	0	1	0	0
---	---	---	---	---	---	---	---	---

# Feature Extraction

Bag of Words representation vectorizes through word counts

Amanda didn't start the fire

	amanda	baby	did	do	fire	not	shark	start	the
Bag of Words	1	0	1	0	1	1	0	1	1
- stop words	1	0	0	0	1	0	0	1	0
+ normalization	0.3	0	0	0	0.3	0	0	0.3	0

# Feature Extraction

One Hot Encoding and TFIDF normalize token frequencies

Baby shark, doo doo doo doo doo doo

	amanda	baby	did	doo	fire	not	shark	start	the
Bag of Words	0	1	0	6	0	0	1	0	0
One Hot Encoding	0	1	0	1	0	0	1	0	0
TFIDF	0	0.05	0	0.4	0	0	0.2	0	0



# Feature Extraction

Transforming text data into numeric features

Non-distributed			Distributed
One Hot Encoding	Bag of Words	TF-IDF	Embeddings
Stop word removal?	Stop word removal	No need for stop word removal	Stop word removal?
None	Document-level normalization	Corpus and document-level normalization	Context and corpus-level normalization

# Word embeddings capture semantic similarities

Statistical modeling through software (e.g. SPSS) or programming language (e.g. **Python**)

**Context**

**Word**

Experience in **Python**, Java or other object-oriented programming languages

**Context**

**Word**

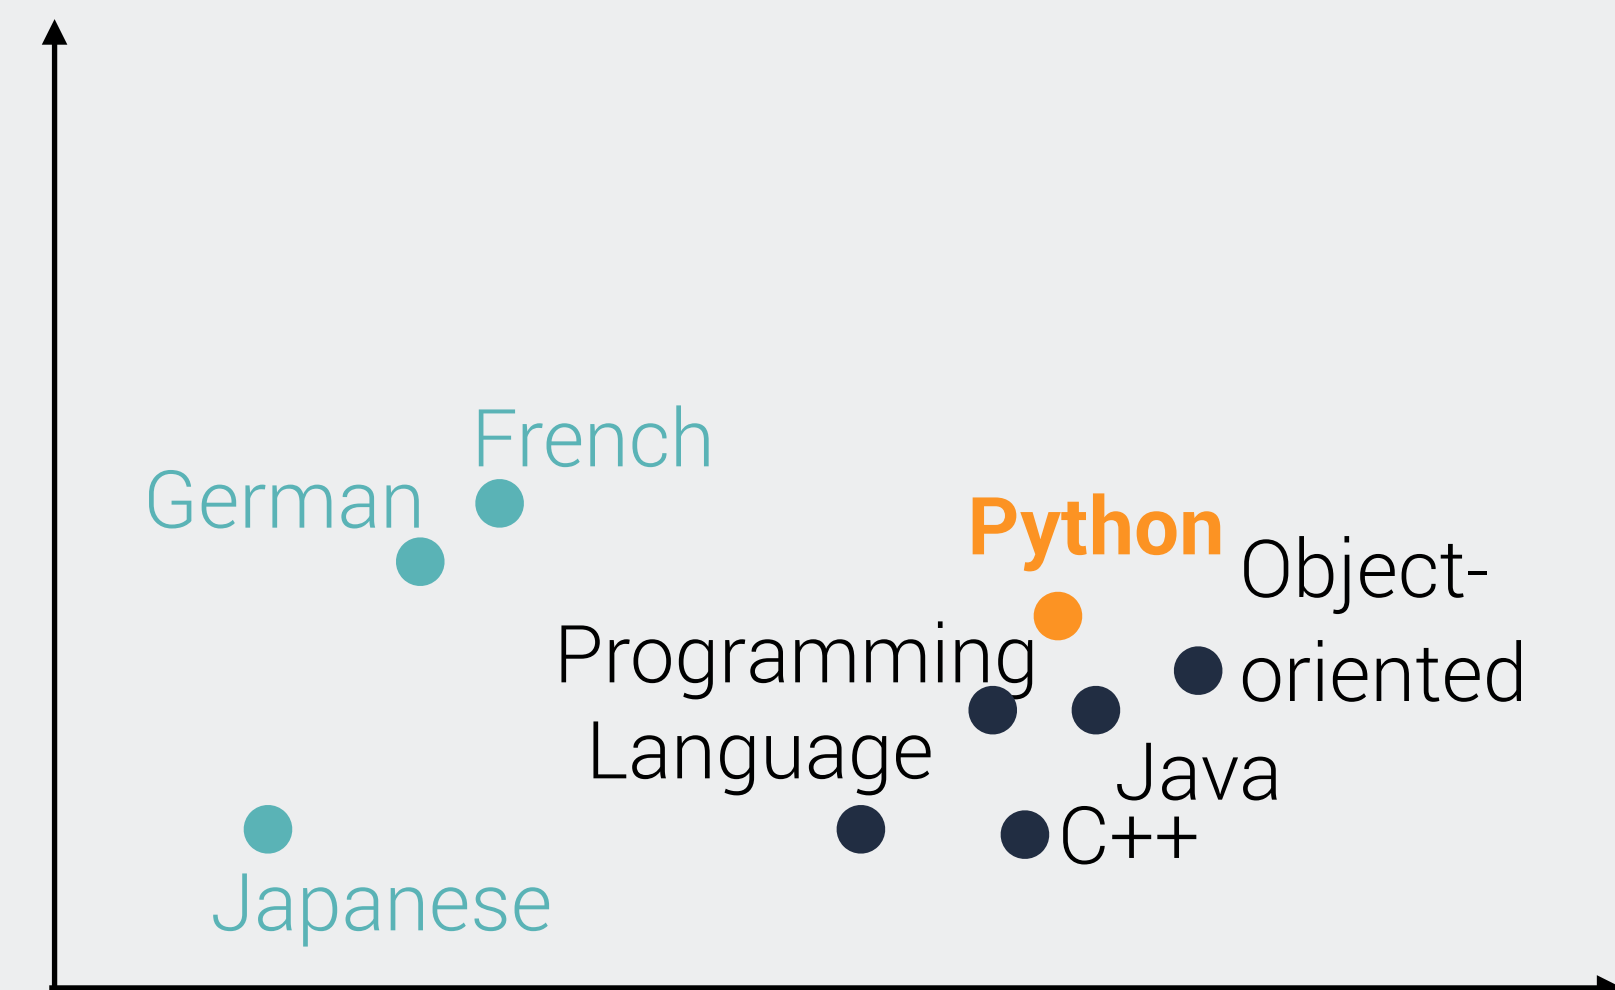
**Context**

Proficiency programming in **Python**, Java or C++.

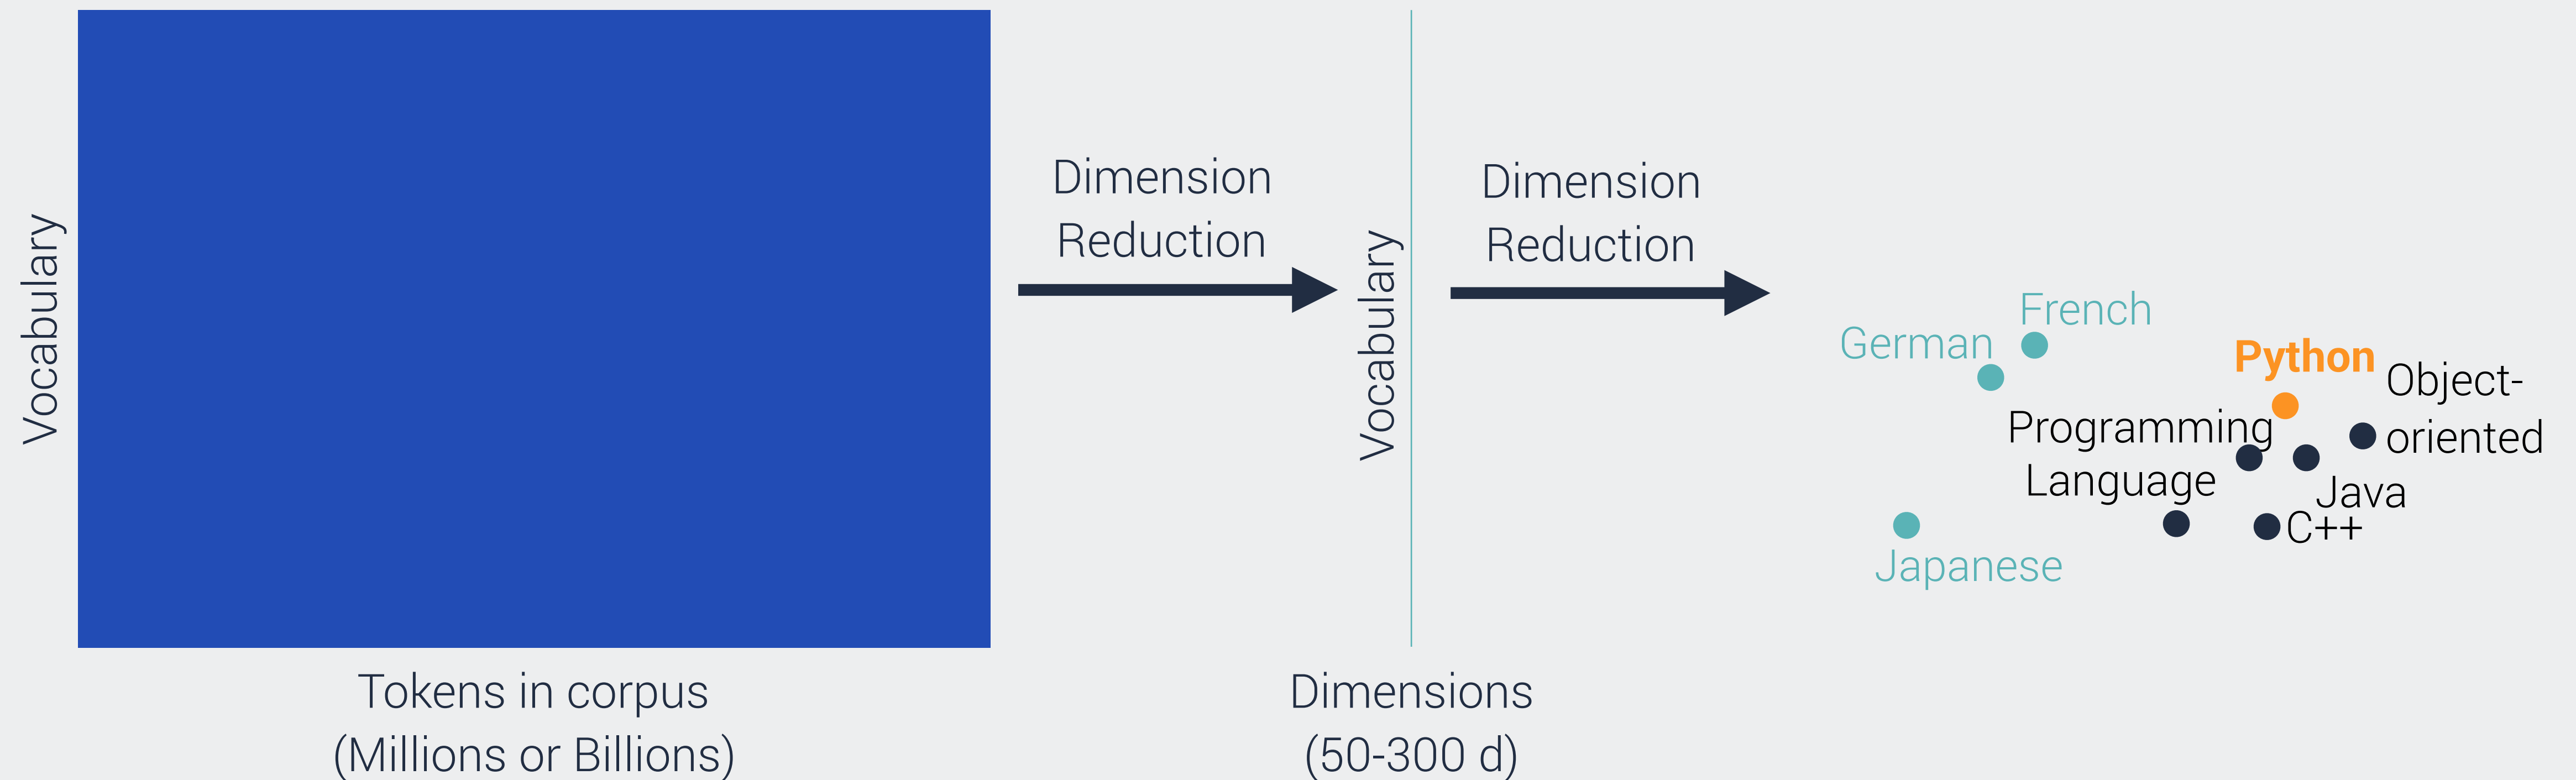
**Context**

**Word**


**Context**



# A simplified representation of word vectors



# A simplified representation of word vectors





# Feature Extraction

Transforming text data into numeric features

Non-distributed			Distributed
One Hot Encoding	Bag of Words	TF-IDF	Embeddings
Stop word removal?	Stop word removal	No need for stop word removal	Stop word removal?
None	Document-level normalization	Corpus and document-level normalization	Context and corpus-level normalization
	All tokens are equidistant		Distance $\propto$ token similarity
	High dimensionality, extremely sparse		Lower dimensionality

# Document categorization

Extracting semantic structure from numeric features

## Topic Modeling

What are the topics that occur in a collection of documents?

Unsupervised Dimension Reduction  
(LDA / LSA)

Every document is a mixture of topics

Every topic is a mixture of words

## Document Classification

Which class does document X belong to?

Supervised ML Algorithms  
Regex (Standard or Artisanal)

Every document belongs to a single class

The presence or absence of a subset of words impacts the classification

# Further Reading

Data Ingestion

**Applied Text Analysis with Python** by  
Benjamin Bengfort, Rebecca Bilbro & Tony  
Ojeda

Data Organization

Data Preprocessing

**Natural Language Processing with Python** by  
Steven Bird, Ewan Klein & Edward Loper

Data Exploration

**Speech and Language Processing** by Dan  
Jurafsky & James Martin

Model Building

**Foundations of Statistical Natural Language  
Processing** by Chris Manning & Hinrich  
Schutze

**Text Mining with R** by  
Julia Silge & David Robinson

# Thank you R Ladies!

**Maryam Jahanshahi Ph.D.**

🐦 @mjahanshahi

in maryam-j

[bit.ly/rladies2019](https://bit.ly/rladies2019)