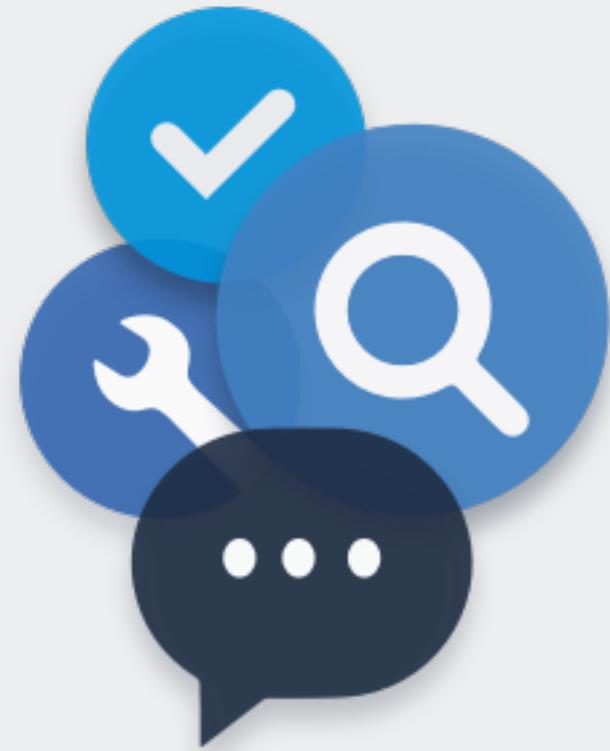


Using Embeddings to Understand the Variance and Evolution of Data Science Skill Sets

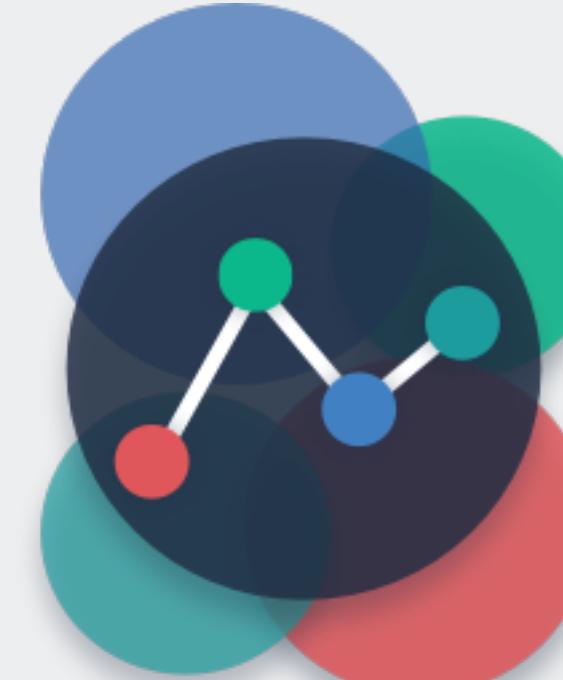
Maryam Jahanshahi Ph.D.
Research Scientist
TapRecruit

TapRecruit uses NLP to understand and organize natural language career content

Smart Editor
for Job Descriptions



Active Pipeline
Health Monitoring



Multifaceted Salary
Estimation





28

Job will perform poorly

This job scores **lower than 95%** of Junior Accounting jobs in Los Angeles, CA



- Add preferred qualifications
- Add more "you" statements
- Perks included
- Equal opportunity statement is included

Neutral



Senior Finance Analyst

TapRecruit - Los Angeles

\$76,300 BETA
\$65,200 \$98,600

TapRecruit is looking for a smart, detail-oriented person to serve as a senior financial analyst. This person will be responsible for supporting the company's FP&A requirements. Responsibilities will include working on TapRecruit Entertainment Group's FP&A model, supporting analysis for long-term planning, tracking key business operational metrics and producing monthly financial/operational reports. The role will require strong organizational skills to help manage the senior managers across the department and evaluate/implement management. This is a dynamic role that serves the finance department of Finance and will routinely interface with TapRecruit's top management.

Language that emphasizes an "intense" or "confusing" environment is known to deter qualified candidates.

Delete

This is an ideal position for an individual who has gained strong accounting firm and now seeks to apply those skills to a fast-growing entrepreneurial company. Strong quantitative and excel financial modeling skills are a must. The ideal candidate must be comfortable in a dynamic start-up environment, will bring energy and passion to everything he/she does, and will not be afraid to roll up his/her sleeves to tackle challenging analytical assignments.

This job is full-time, based in Los Angeles. We offer competitive compensation and stock option program.

Responsibilities:

- Provide financial support to News group and overall FP&A / corporate finance department

Language matters in job descriptions

**Same title,
Different job**

Finance Manager

Kraft Foods

Junior (3 Years)
No Managerial Experience

Finance Manager

Kraft Foods

Senior (6-8 Years)
Division Level Controller
Strategic Finance Role
MBA / CPA

Same Title

Required Experience
Required Responsibility
Preferred Skill
Required Education

**Different title,
Same job**

Performance Marketing Manager

PocketGems

Mid-Level
Quantitative Focus
Expertise iBanking
Data Analysis Tools (SQL)
Consulting Experience Preferred
MBA Preferred

Senior Analyst, Customer Strategy

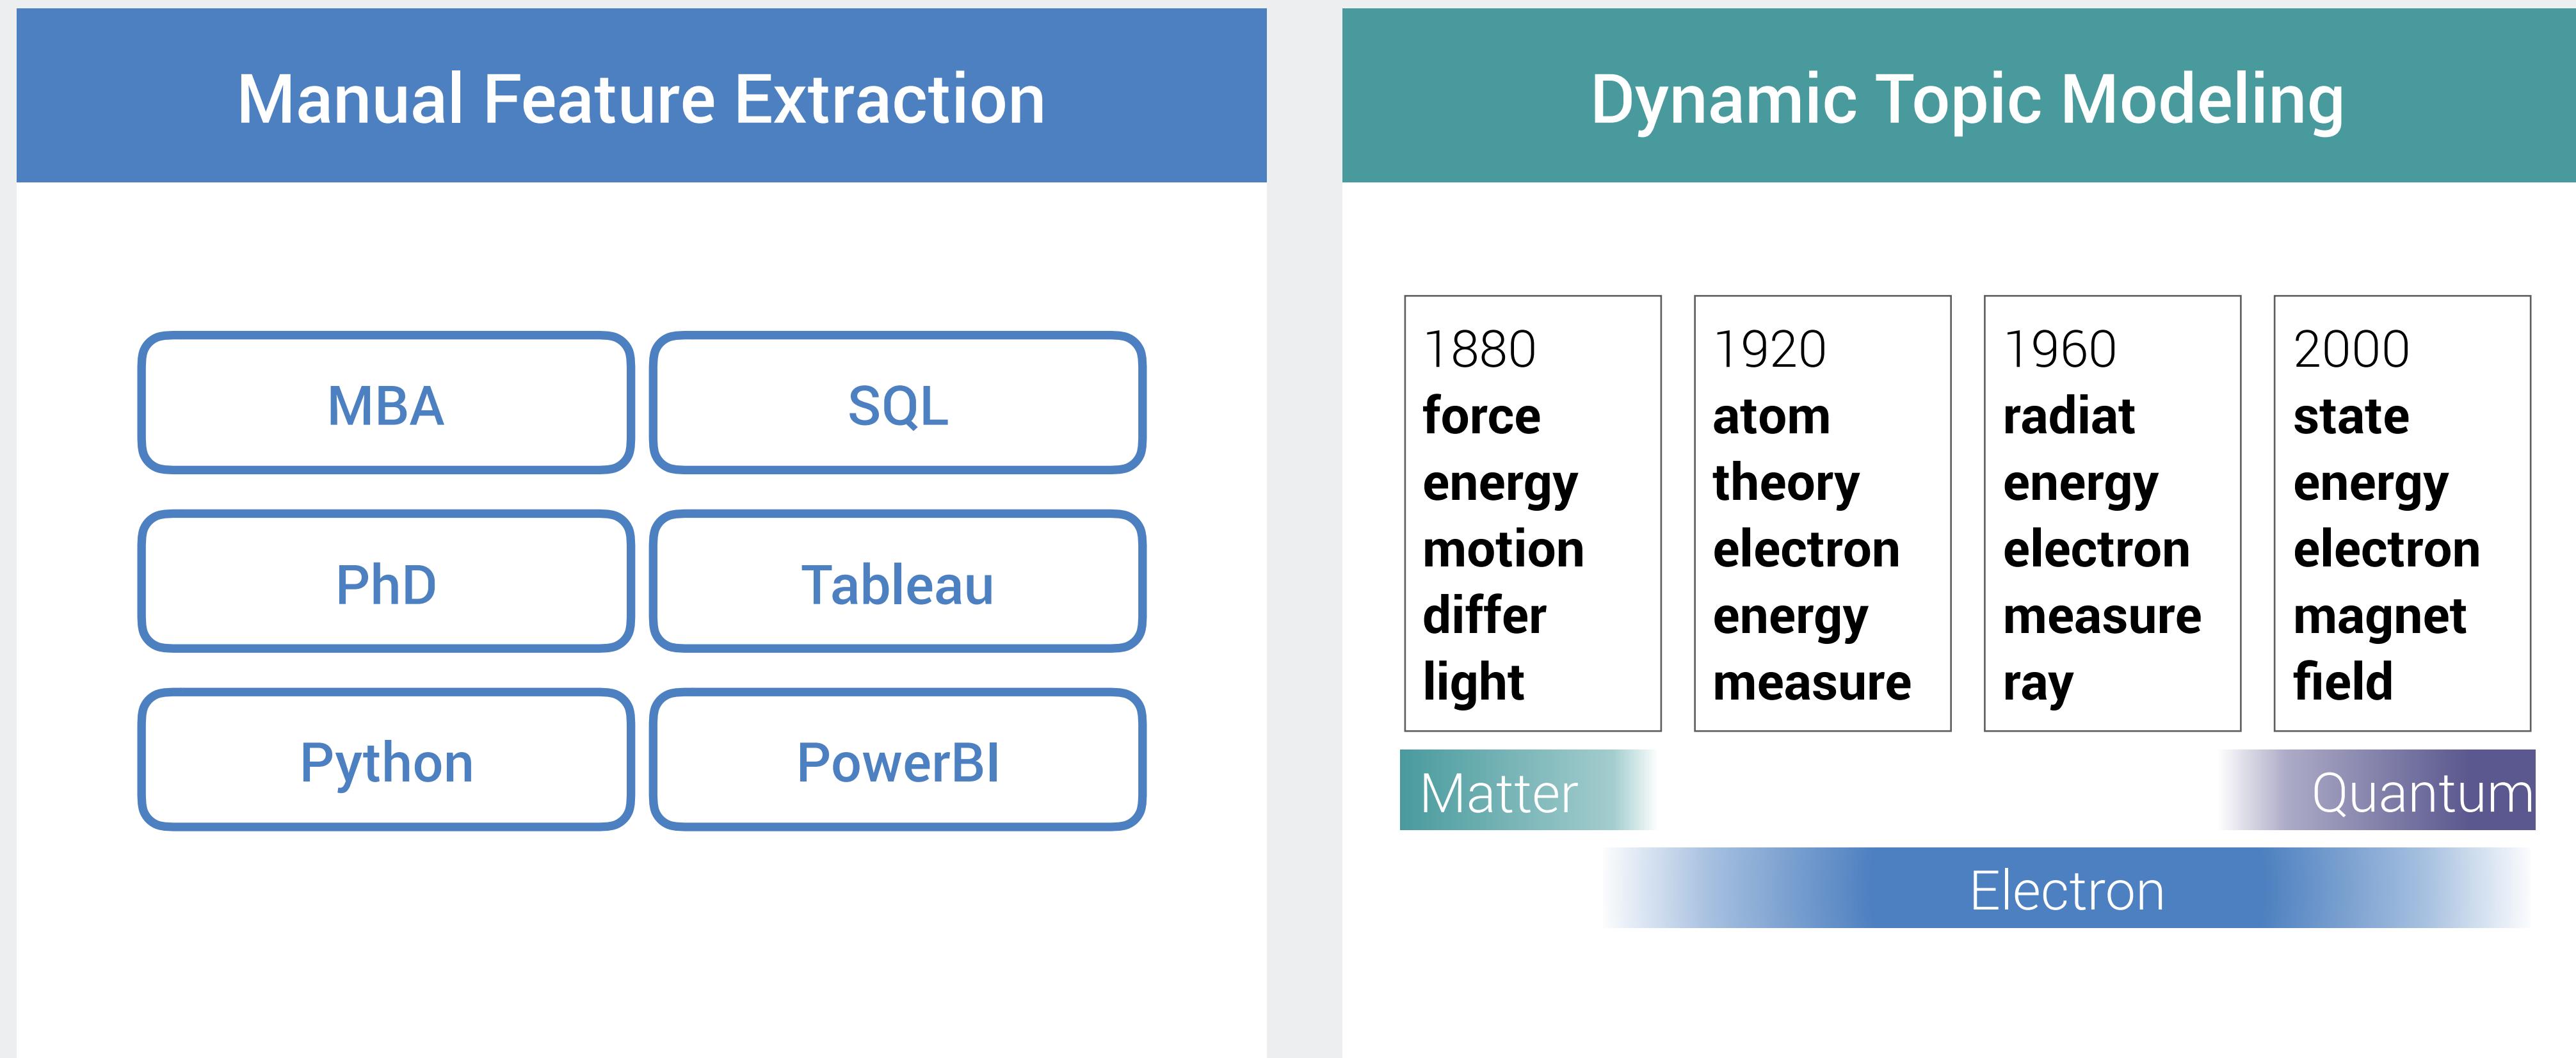
The Gap

Mid-Level
Quantitative Focus
Expertise iBanking
Relational Database Experience
External Consulting Experience Preferred
BA degree in business, finance, MBA
Preferred

Required Experience
Required Skills
Required Experience
Required Skills
Preferred Experience
Required and Preferred Education

How have data science skills
changed over time?

How have data science skills changed over time?



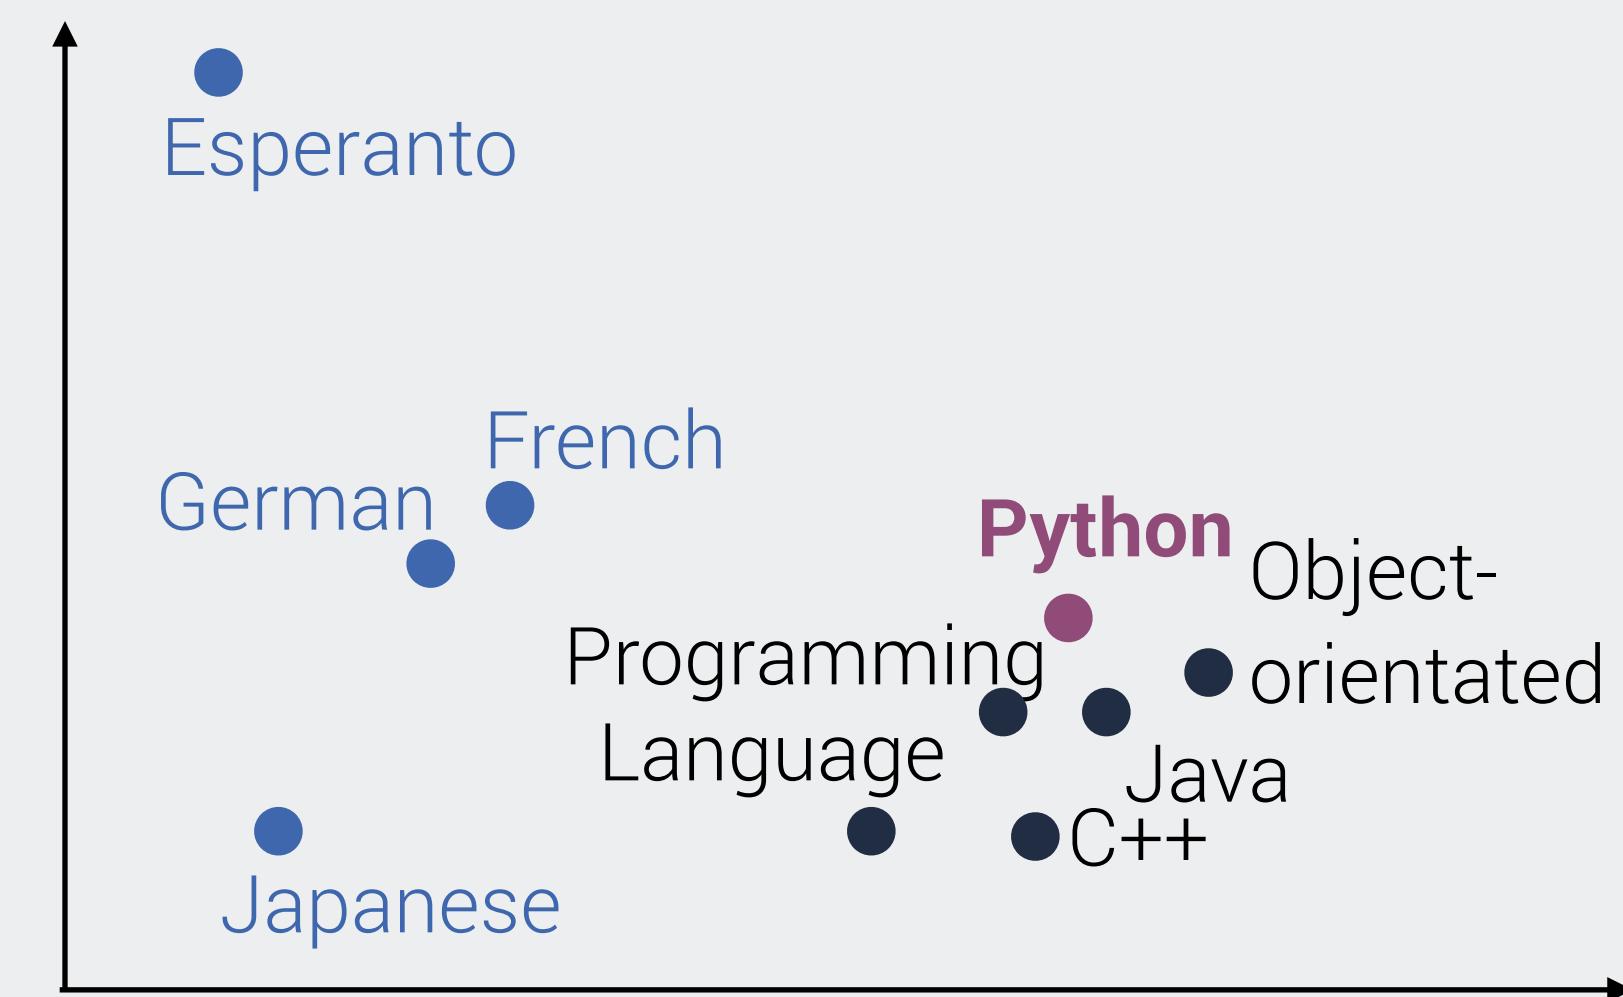
Adapted from Blei and Lafferty, [ICML 2006](#).

Word embeddings capture semantic similarities

Statistical modeling through software (e.g. SPSS) or programming language (e.g. **Python**)

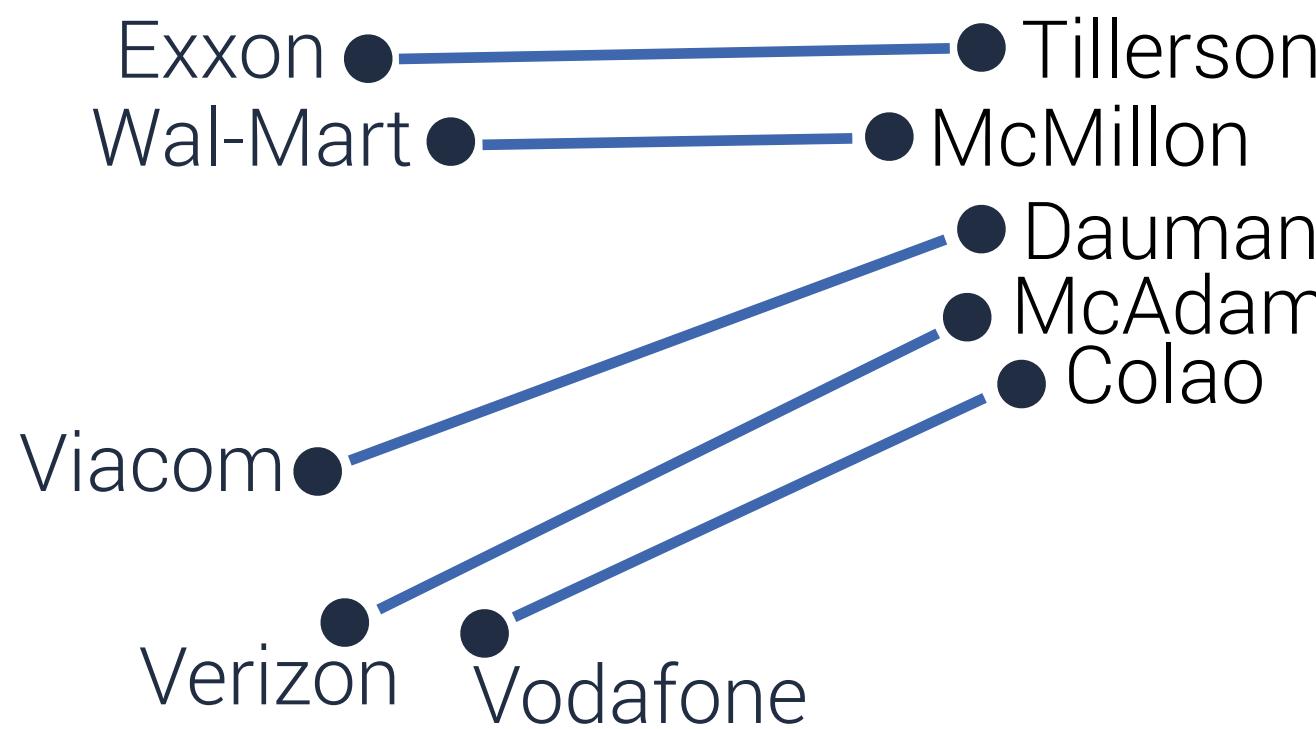
Experience in **Python**, Java or other object-oriented programming languages

Proficiency programming in **Python**, Java or C++.

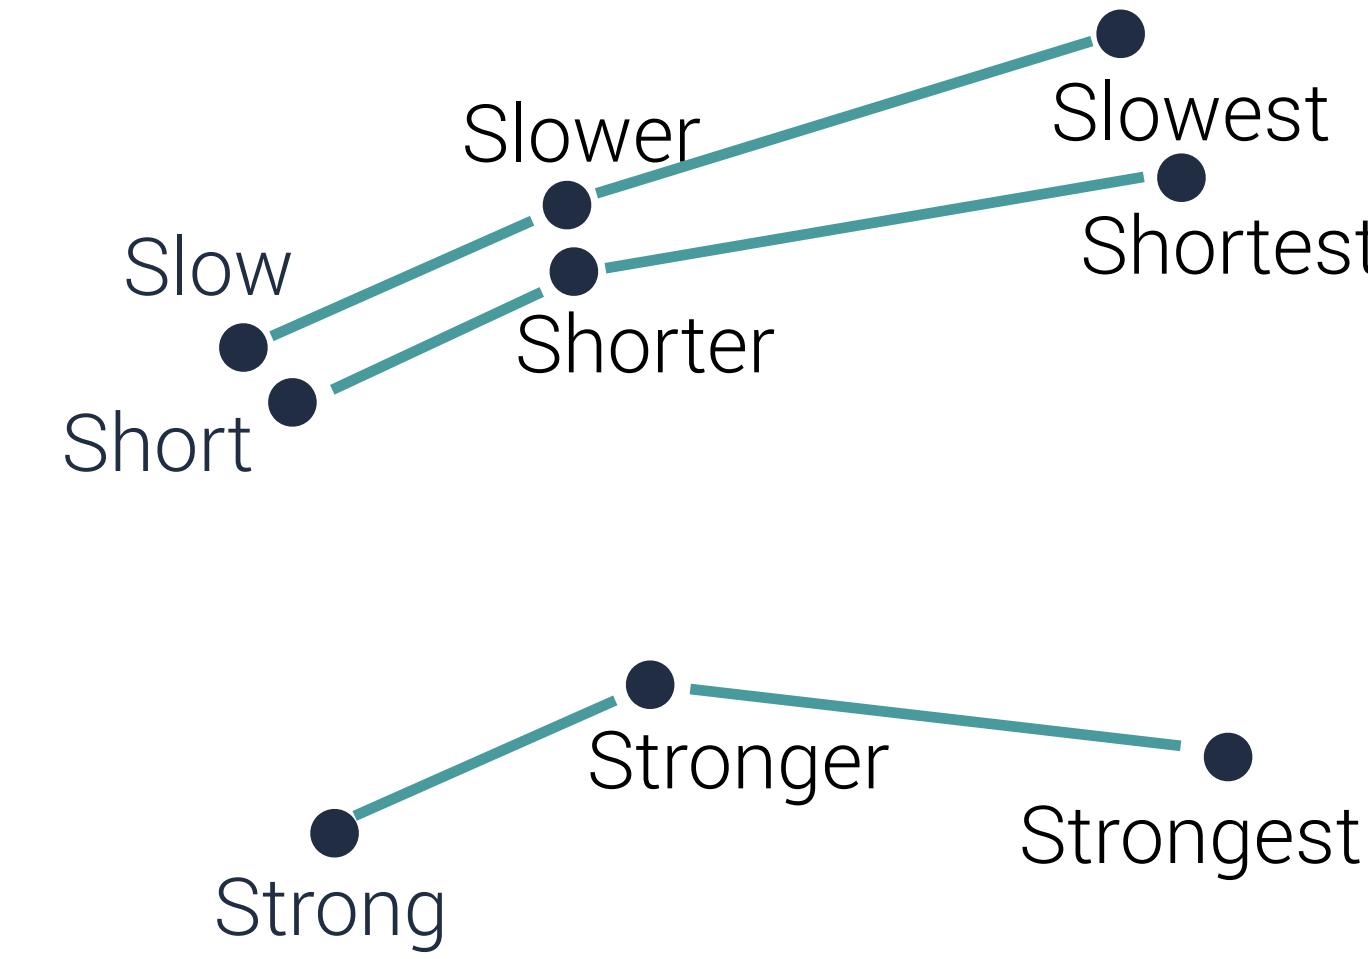


Embeddings capture entity relationships

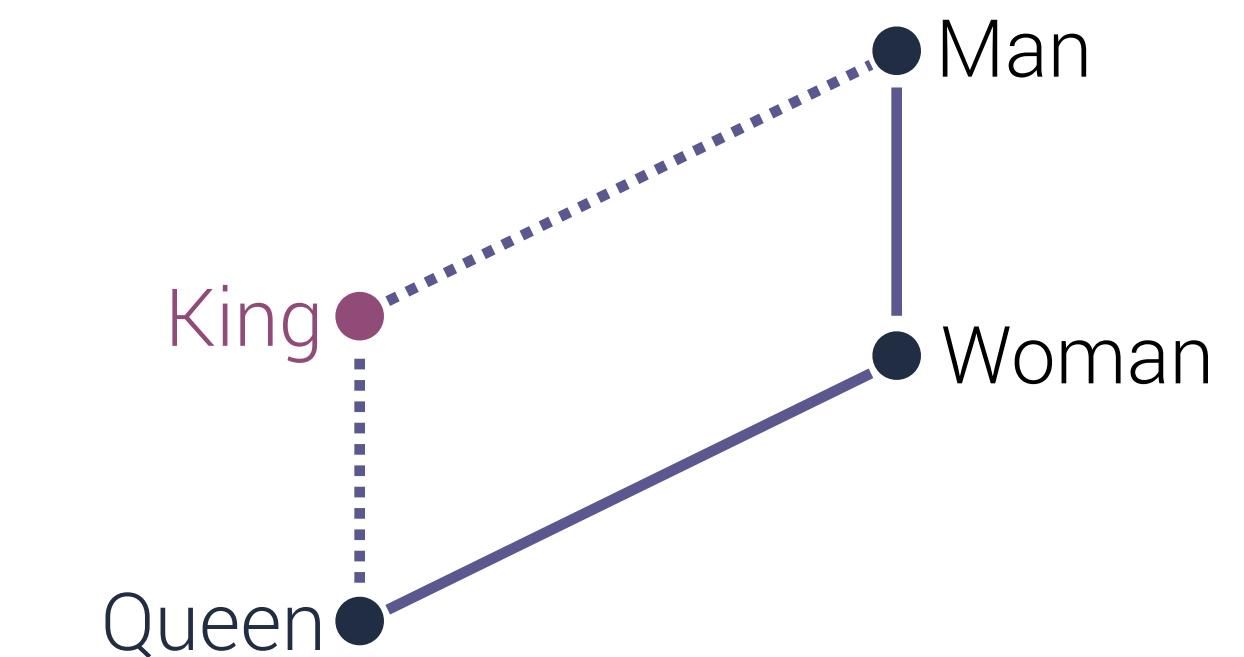
Hierarchies



Comparatives and Superlatives



Woman :: Queen as Man :: ?



Adapted from [Stanford NLP GLoVe Project](#)

Pretrained embeddings facilitate fast prototyping

Corpus Generation	Corpus Tokens	Twitter 27 B	Common Crawl 42-840 B	GoogleNews 100 B	Wikipedia 6 B
Corpus Processing	Vocabulary Size	1.2 M	1.9-2.2 M	3 M	400 k
Language Model Generation	Algorithm Vector Length	GLoVE 25 - 200 d	GLoVE 300 d	word2vec 300 d	GLoVE 50 - 300 d
Language Model Tuning					
Final Application					

Problems with pretrained embedding models

Casing

Abbreviations vs Words
e.g. IT vs it

Out of Vocabulary Words

Domain Specific Words & Acronyms

Polysemy

Words with multiple meanings
e.g. drive (a car) vs drive (results)
e.g. Chef (the job) vs Chef (the language)

Multi-word Expressions

Phrases that have new meanings
e.g. Front-end vs front + end

Tools for developing custom language models

Corpus Processing

Tokenization, POS tagging, Sentence Segmentation, Dependency Parsing



Language Modeling

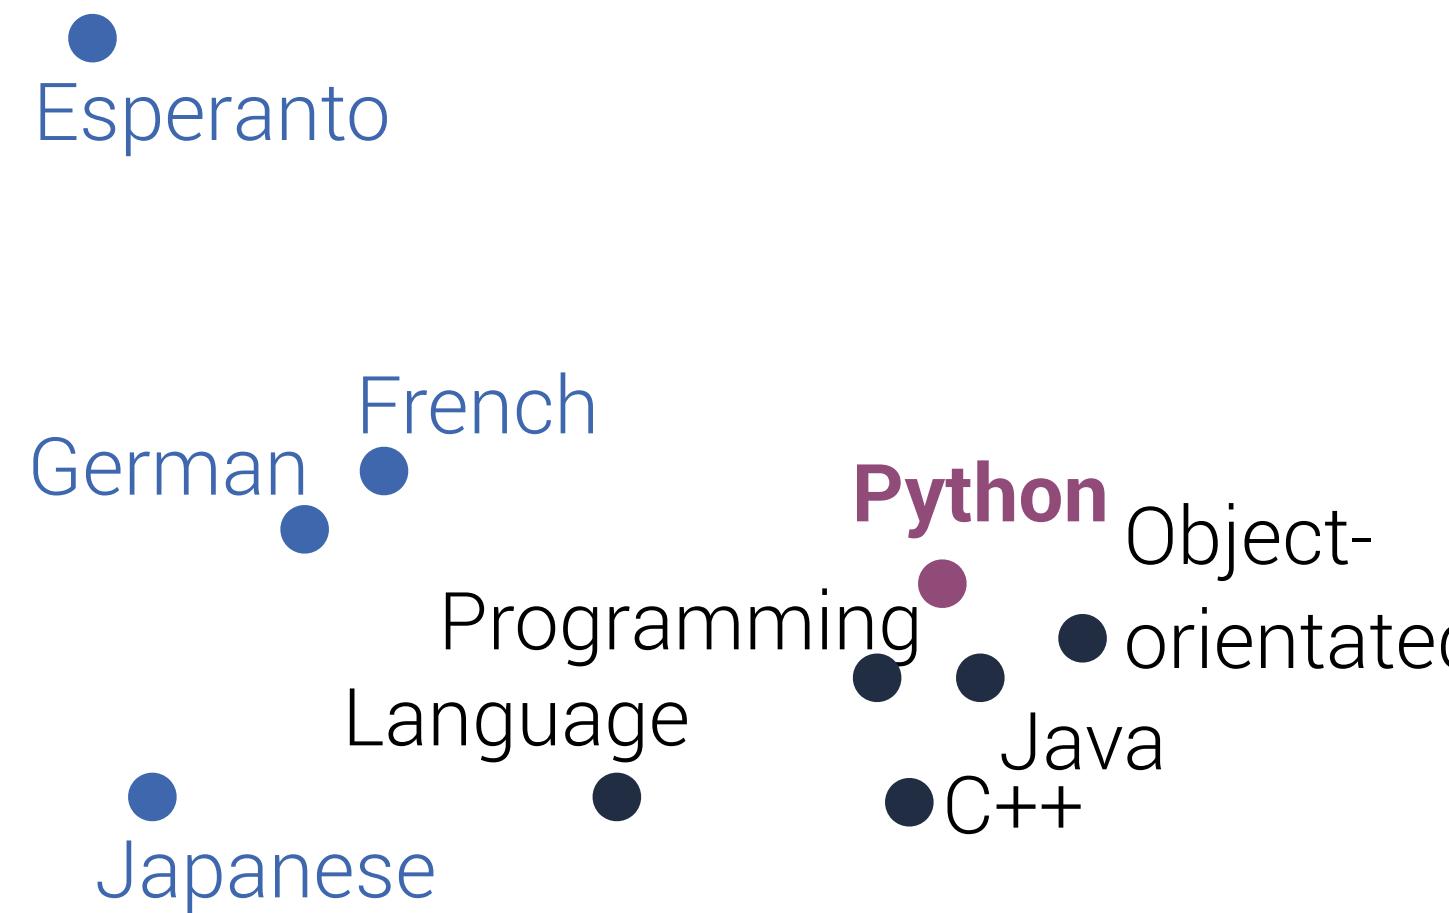
Different word embedding models (GLoVE, word2vec, fastText)



Windows capture semantic similarity vs relatedness

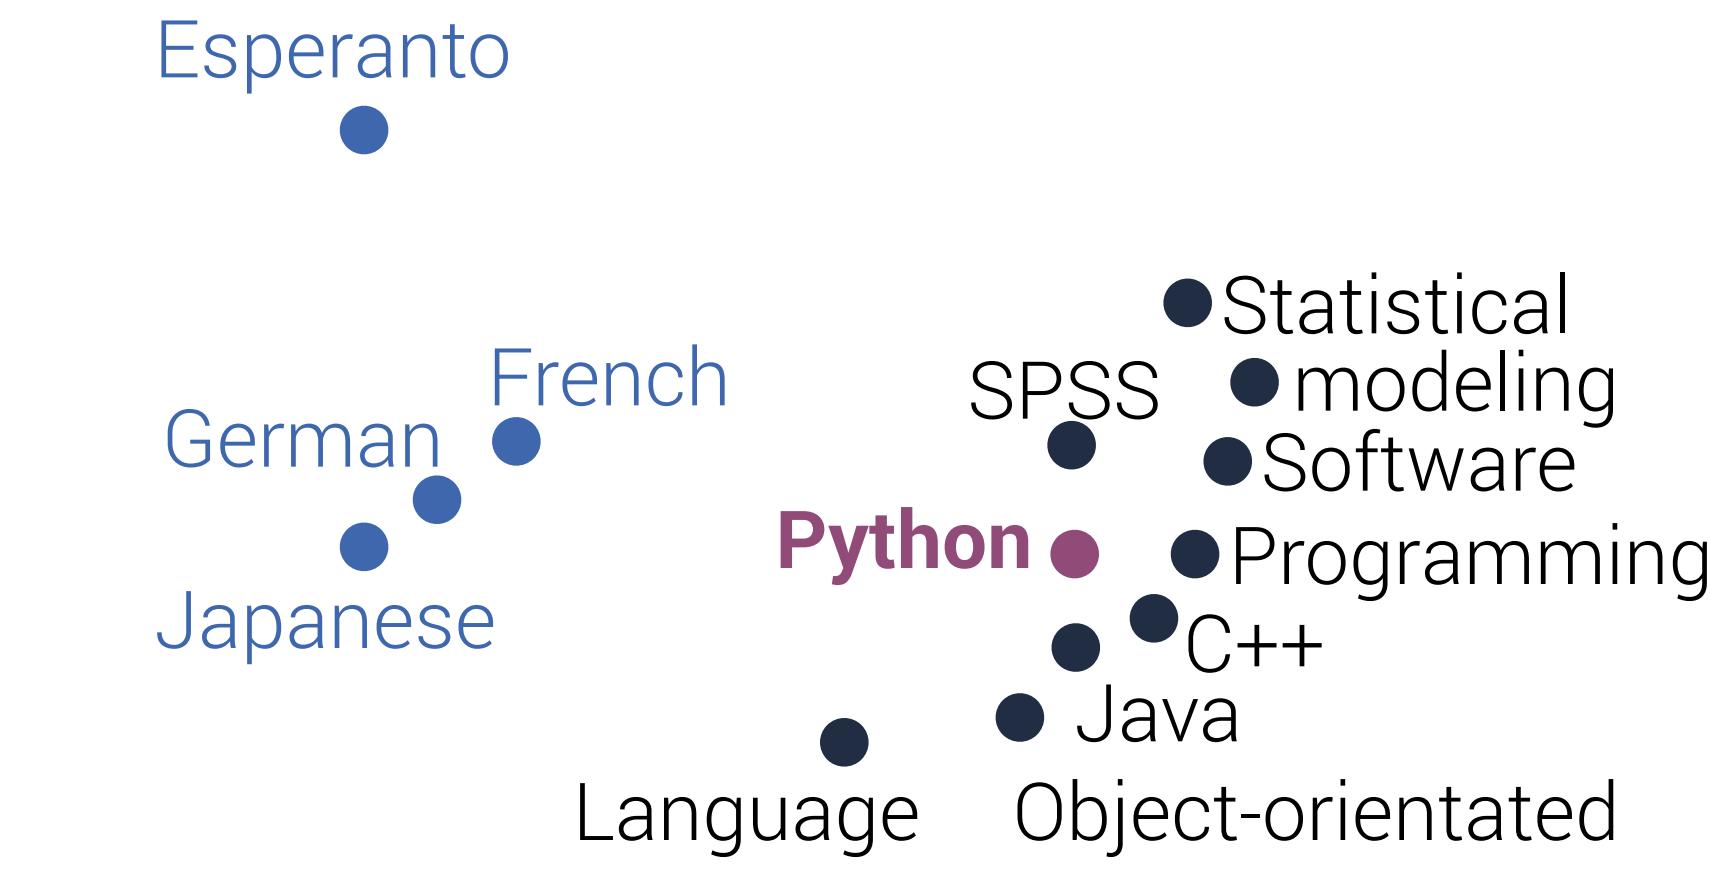
Small Window Size

Captures Semantic similarity, Substitutes and Word-level differences

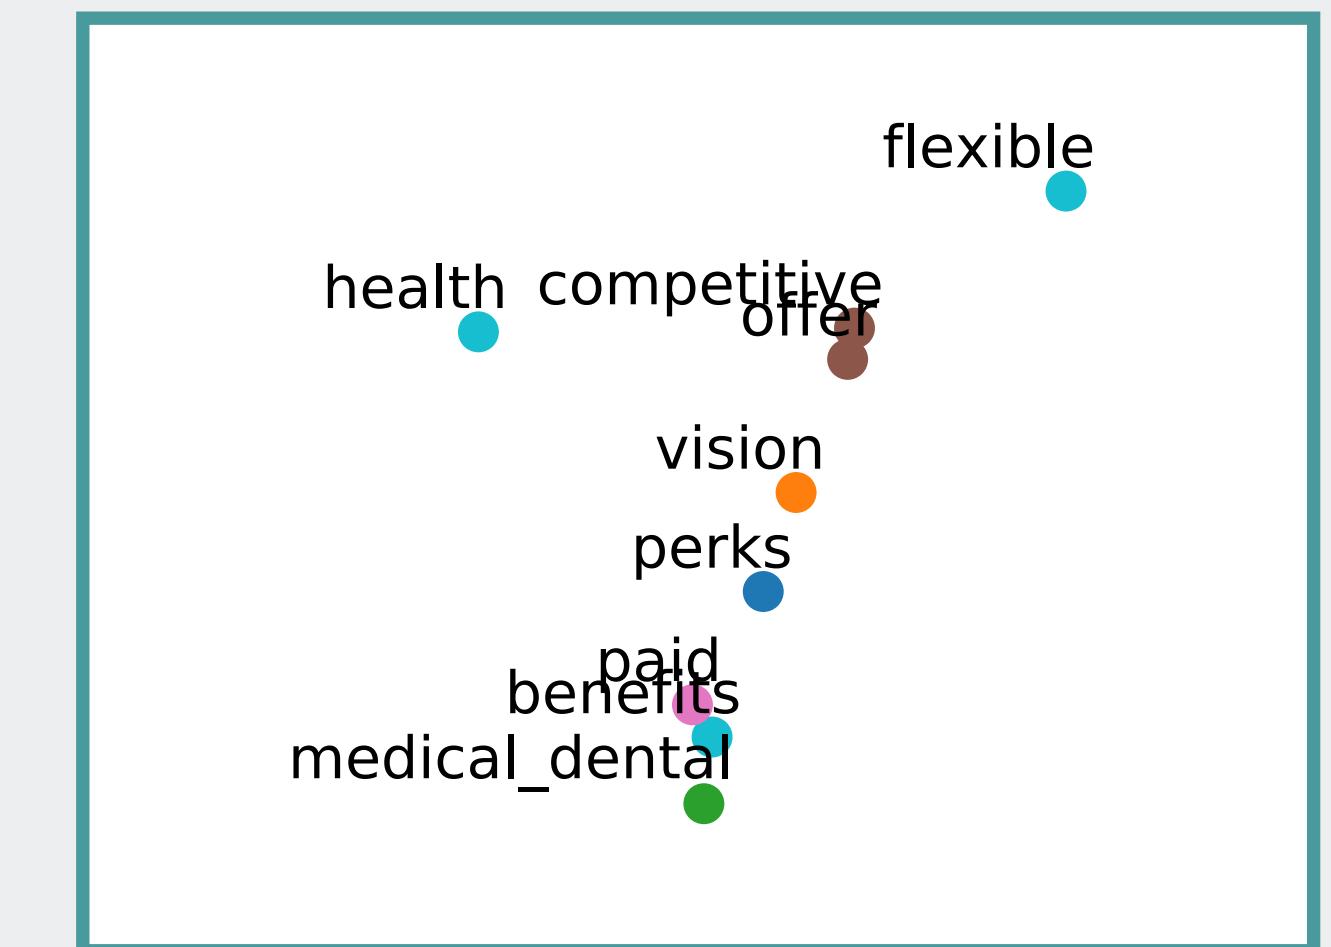
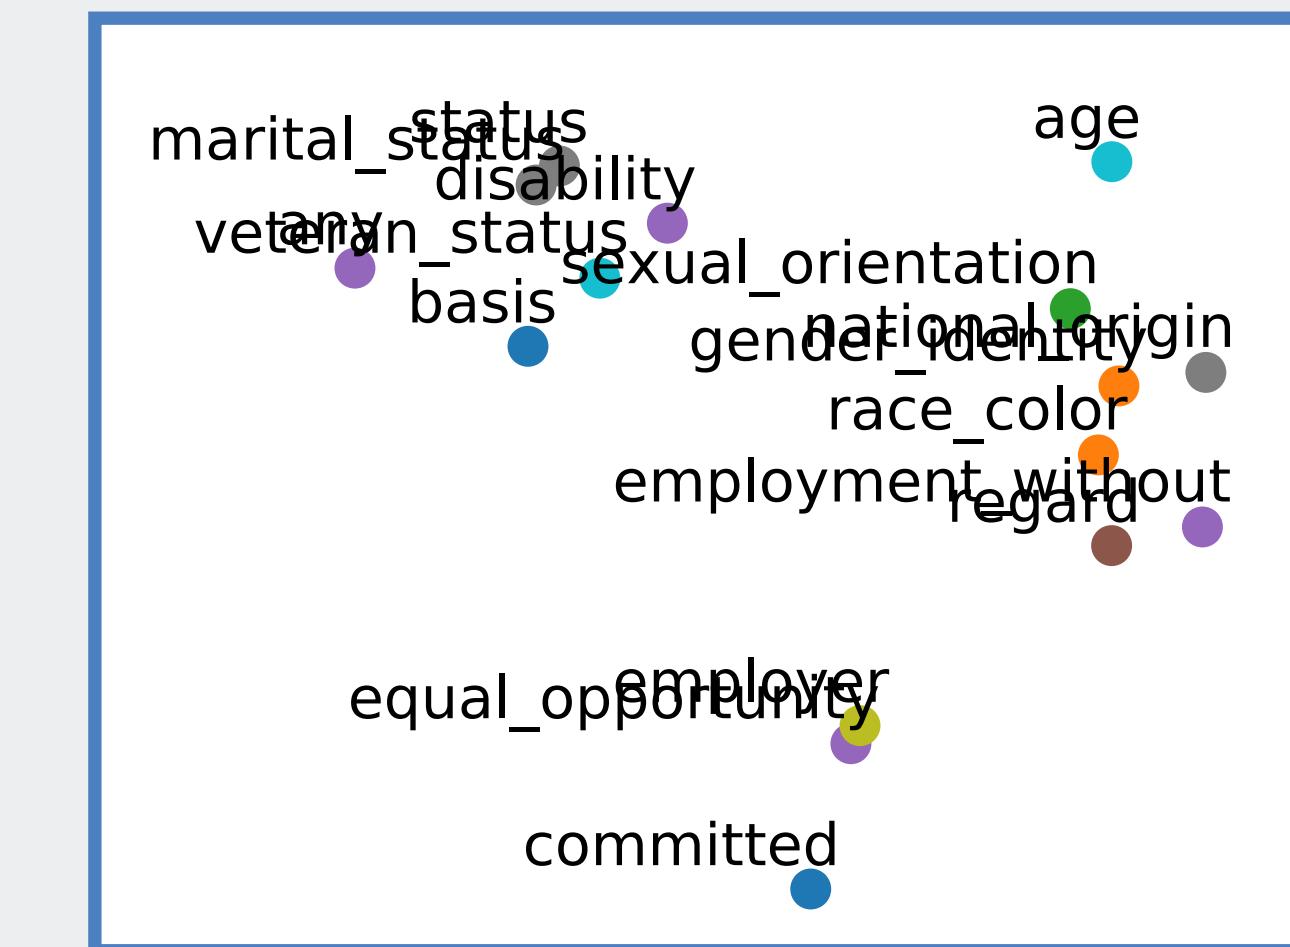
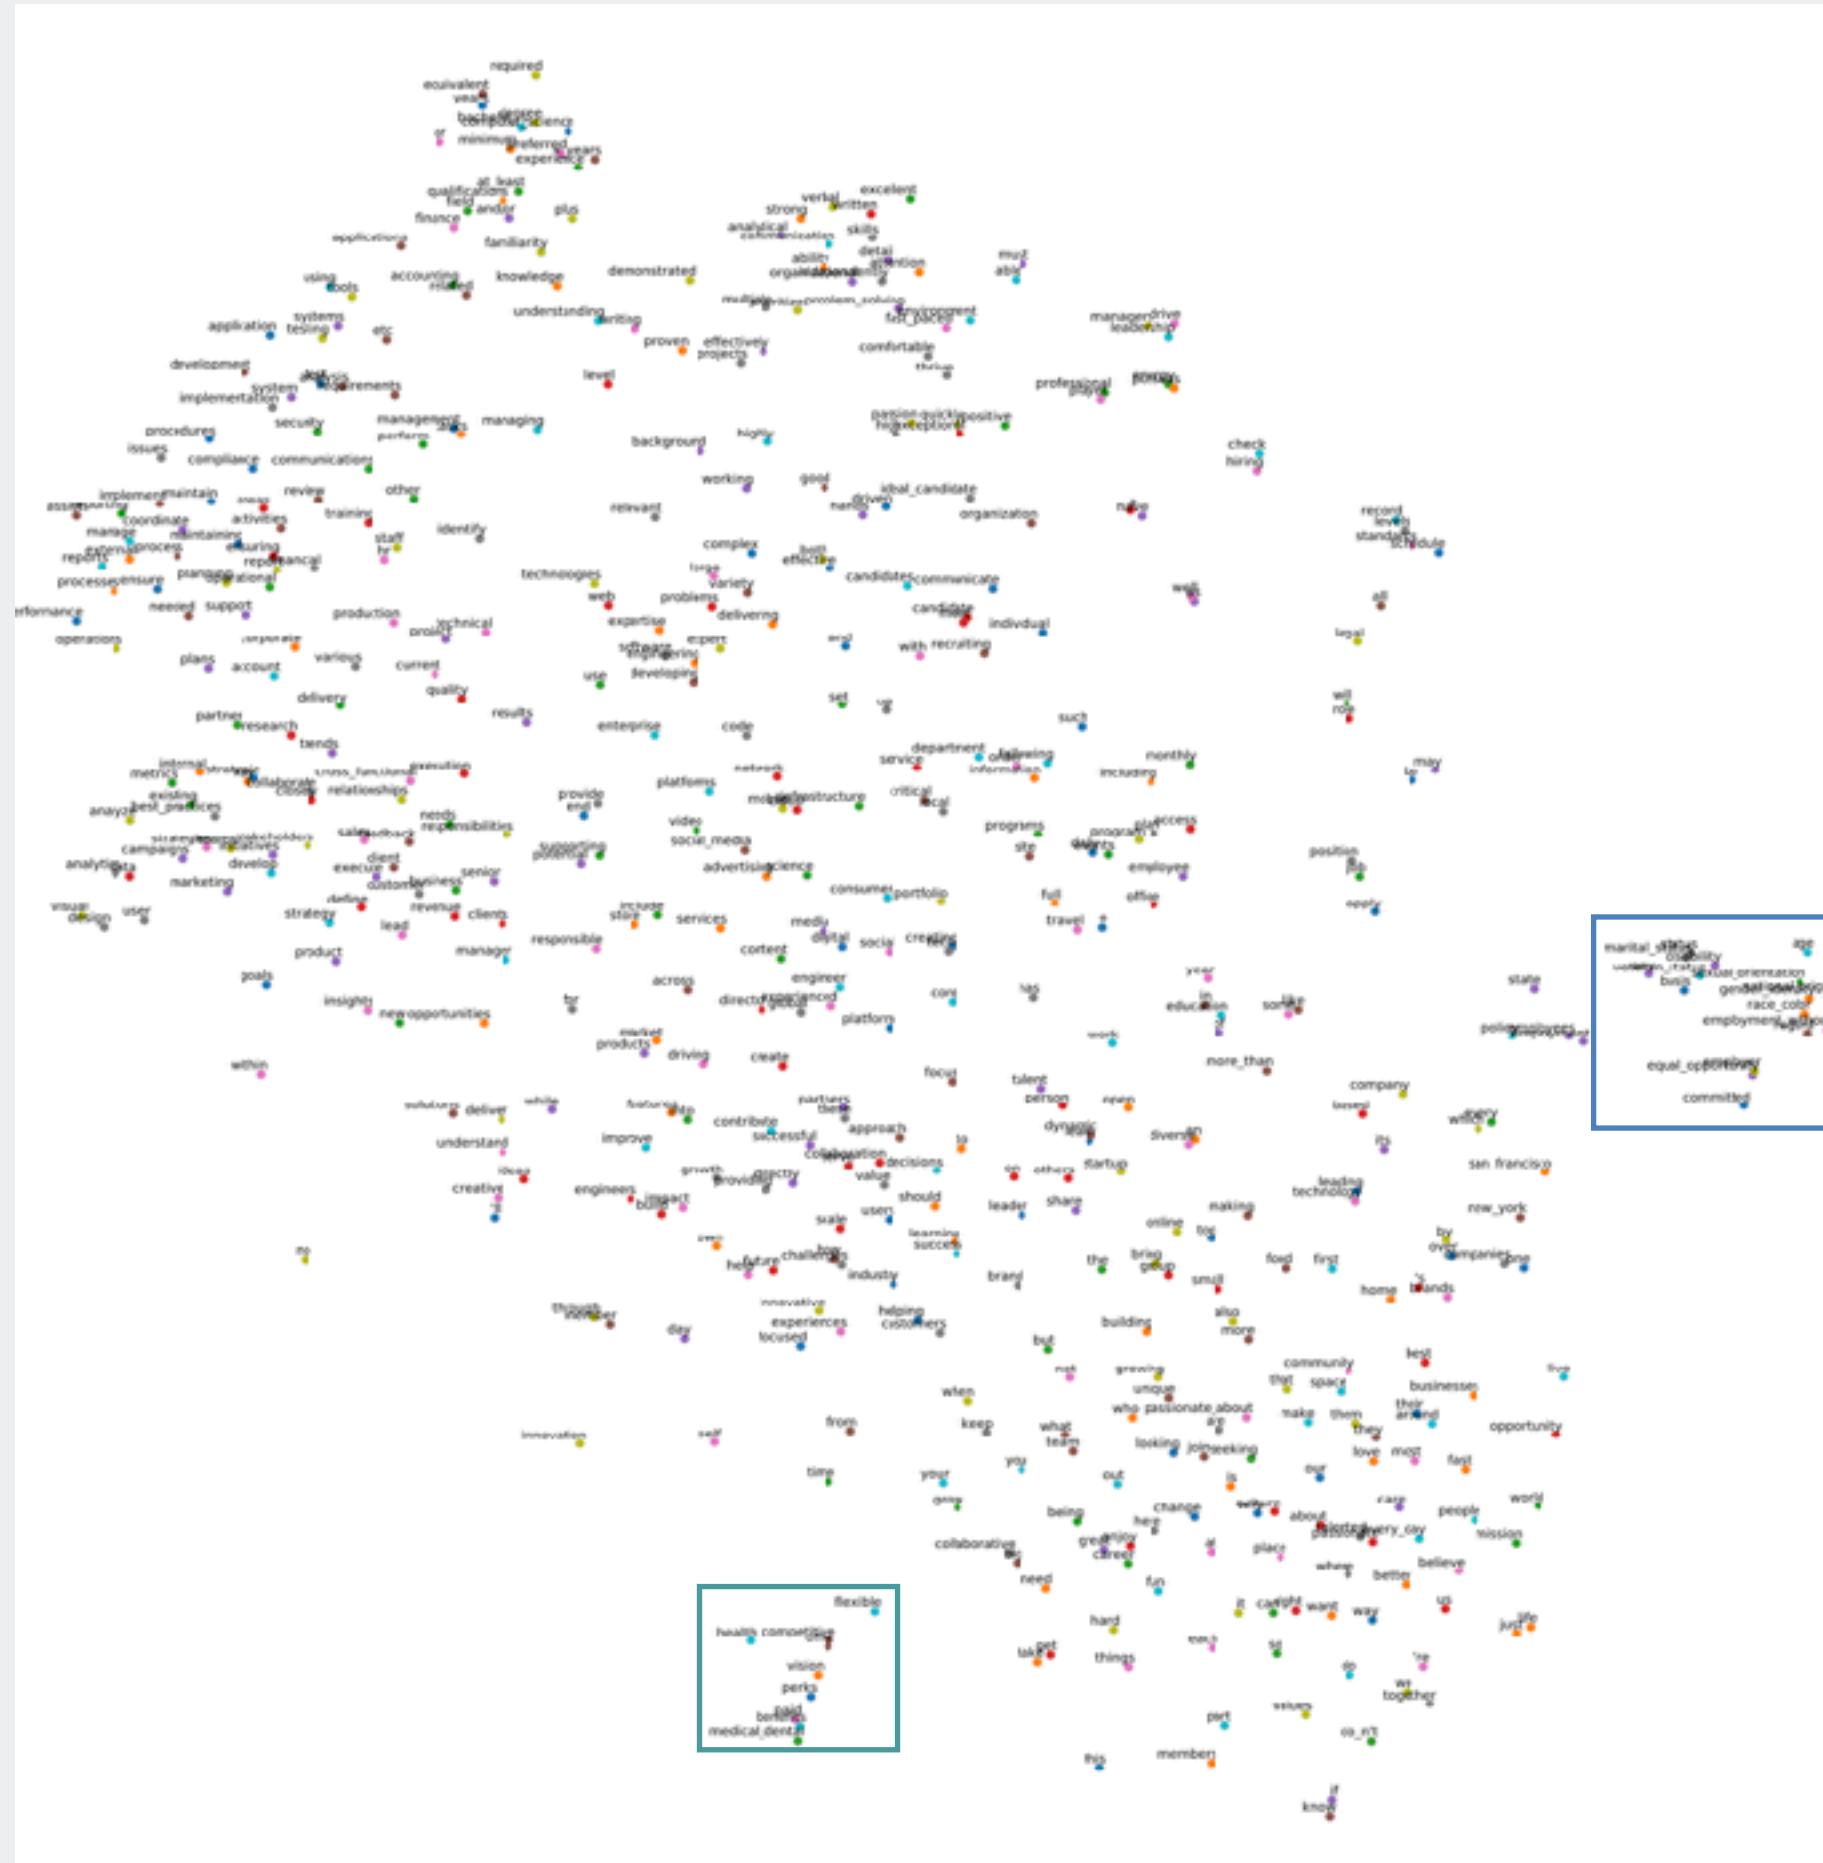


Large Window Size

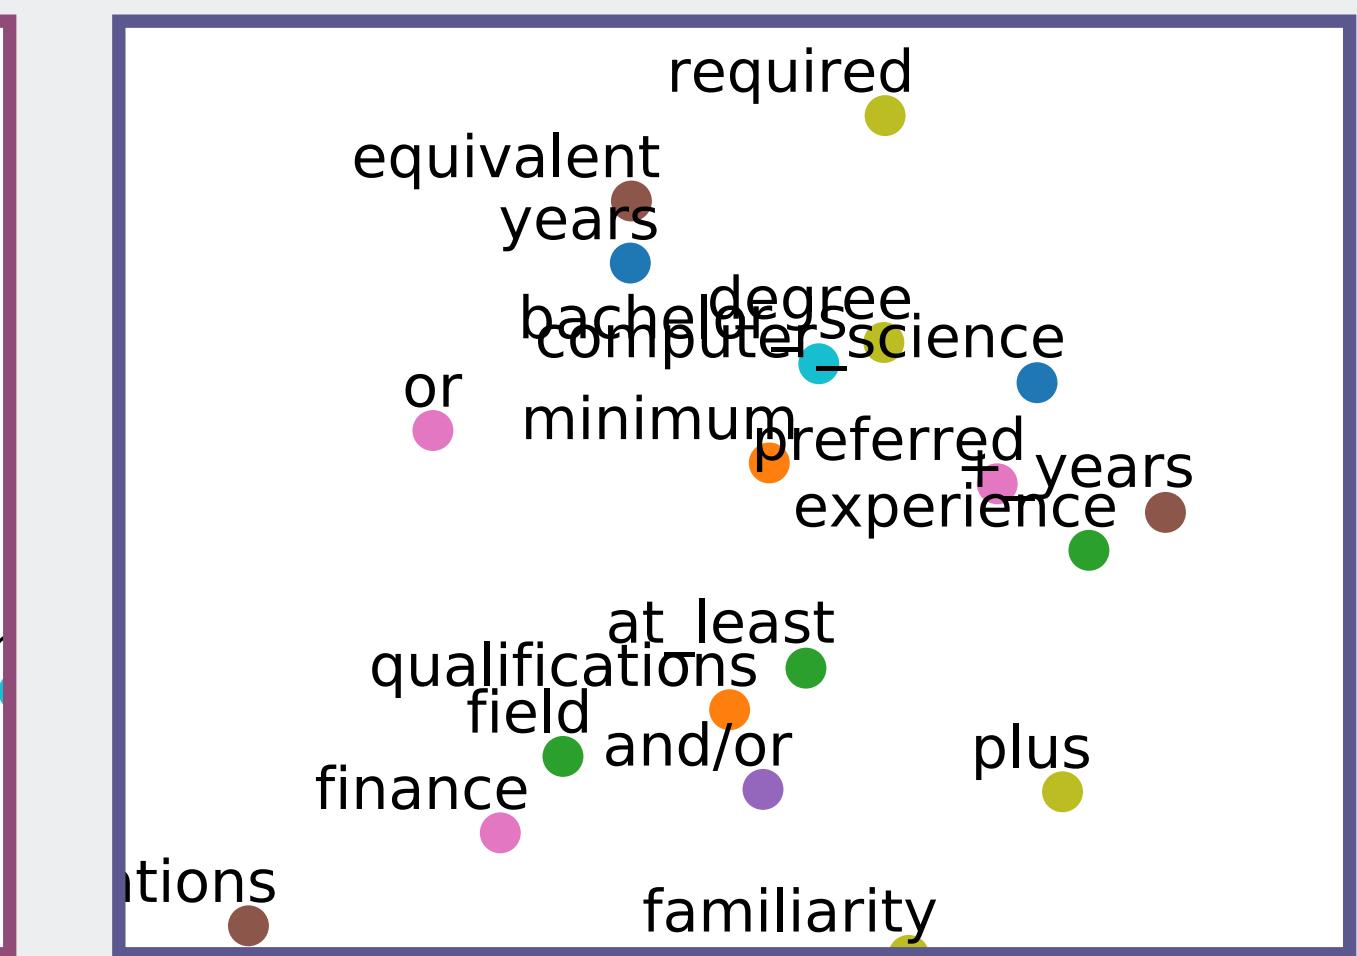
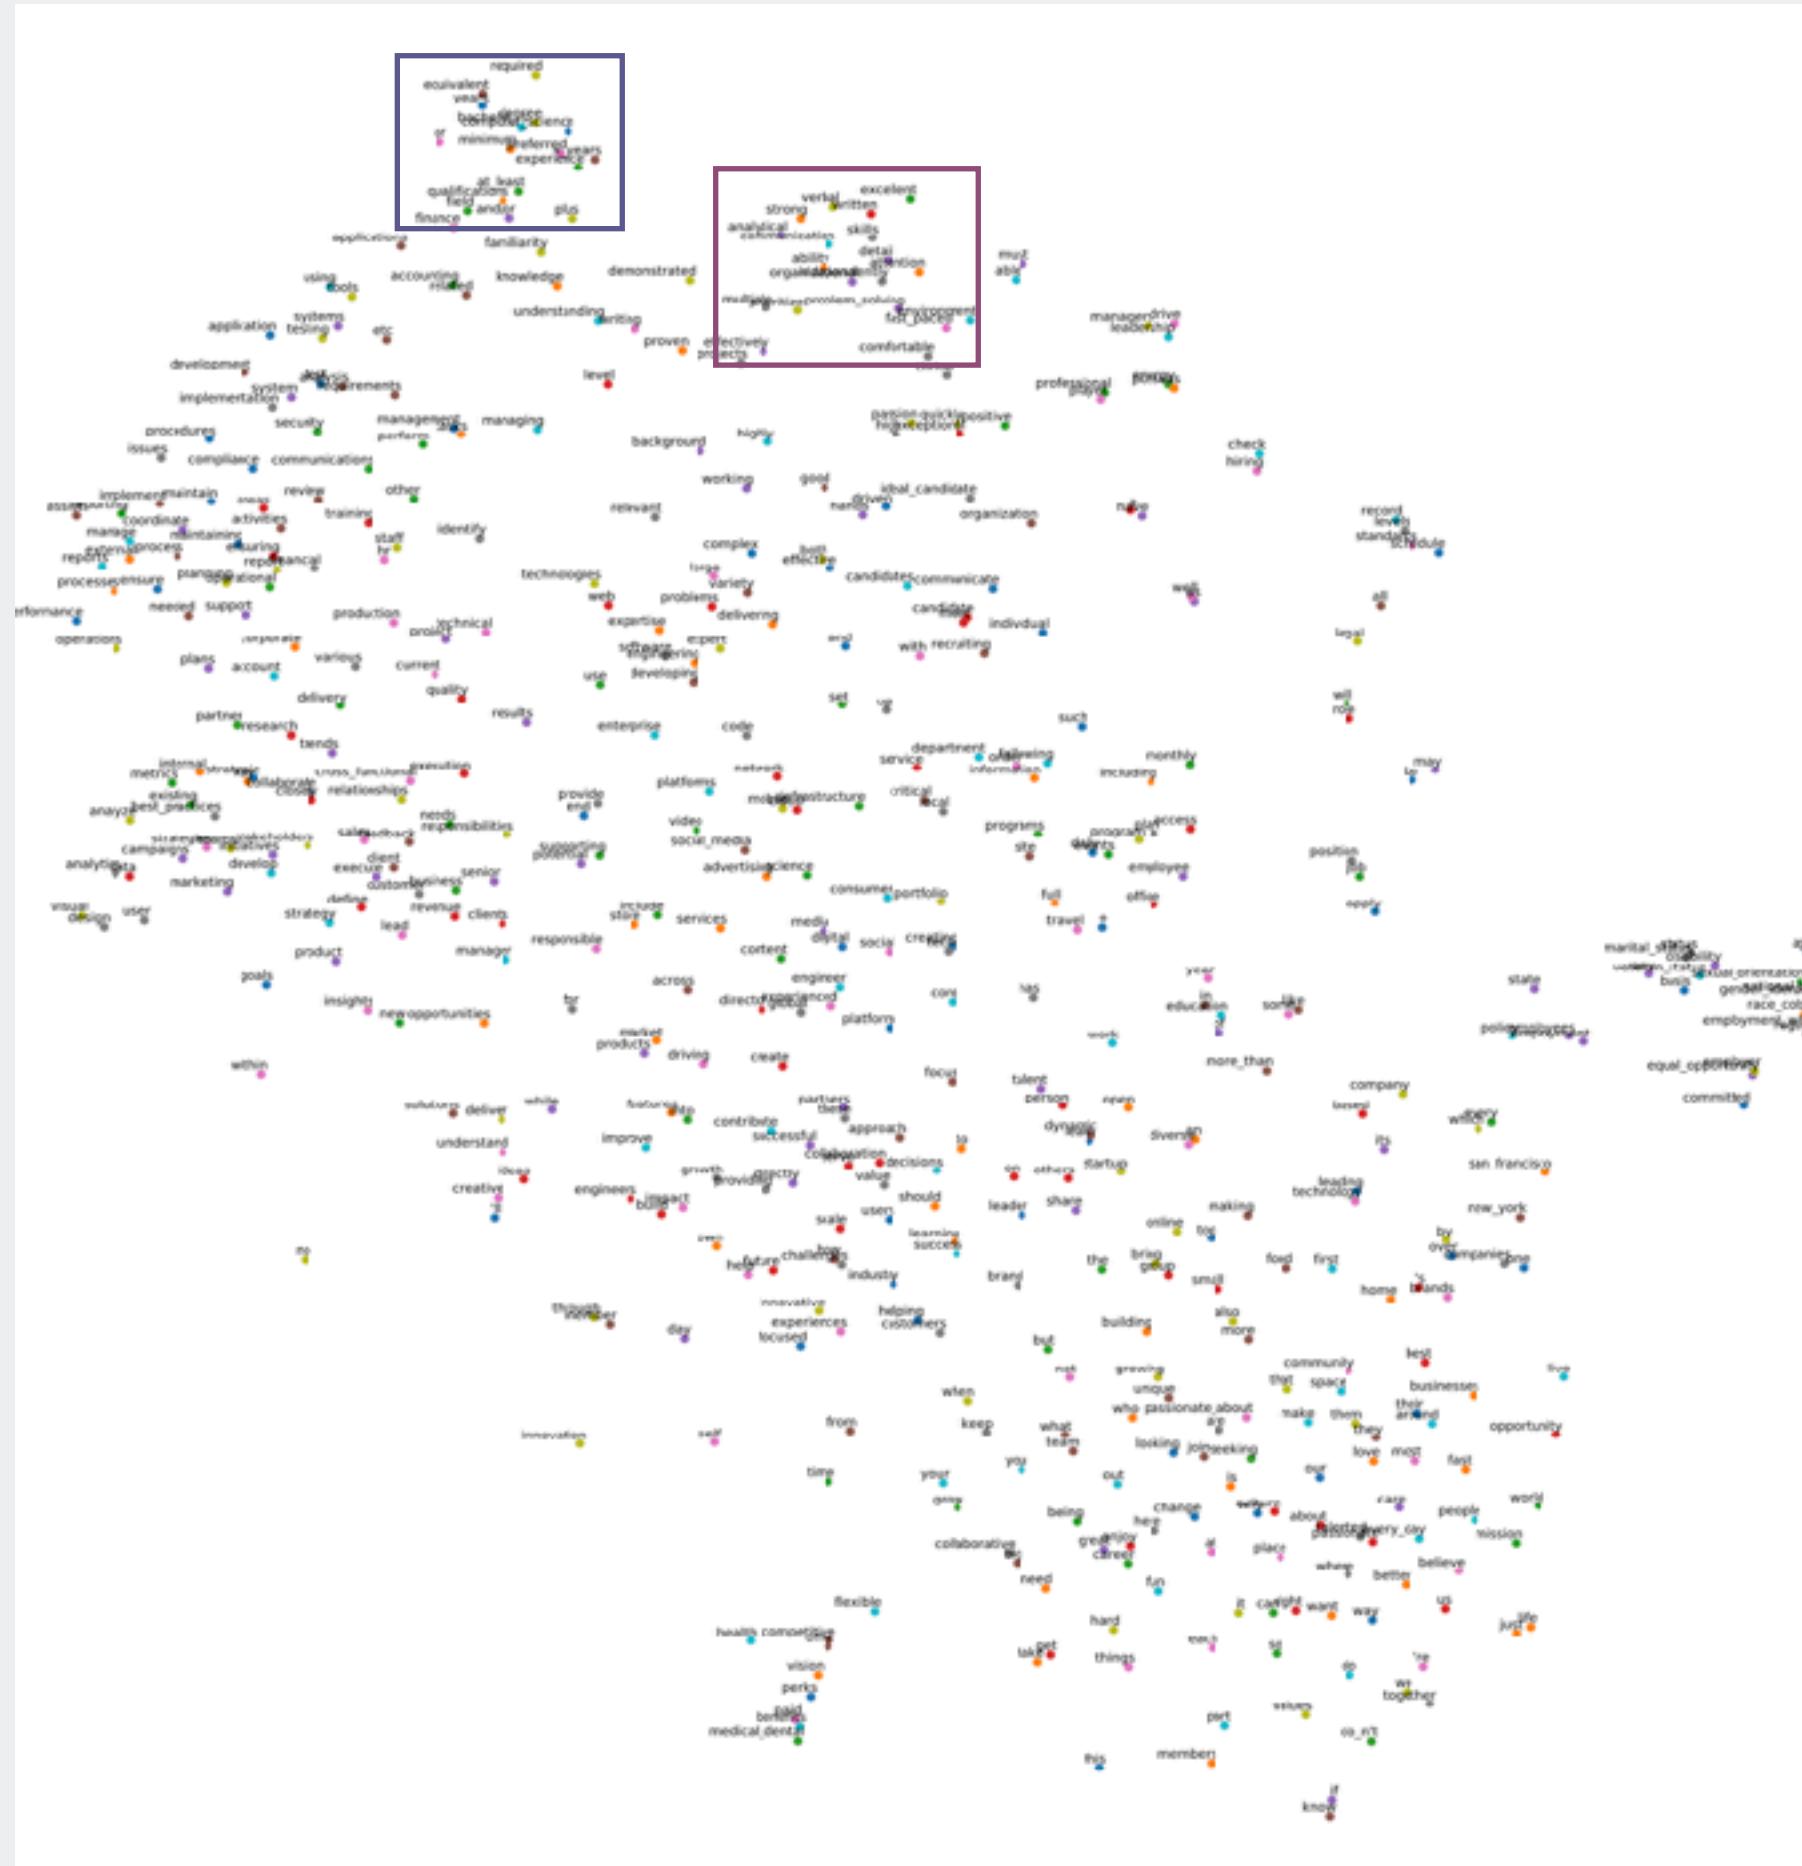
Captures Semantic relatedness, Alternatives and Domain-level differences



Custom embeddings identified equal opportunity and perks language

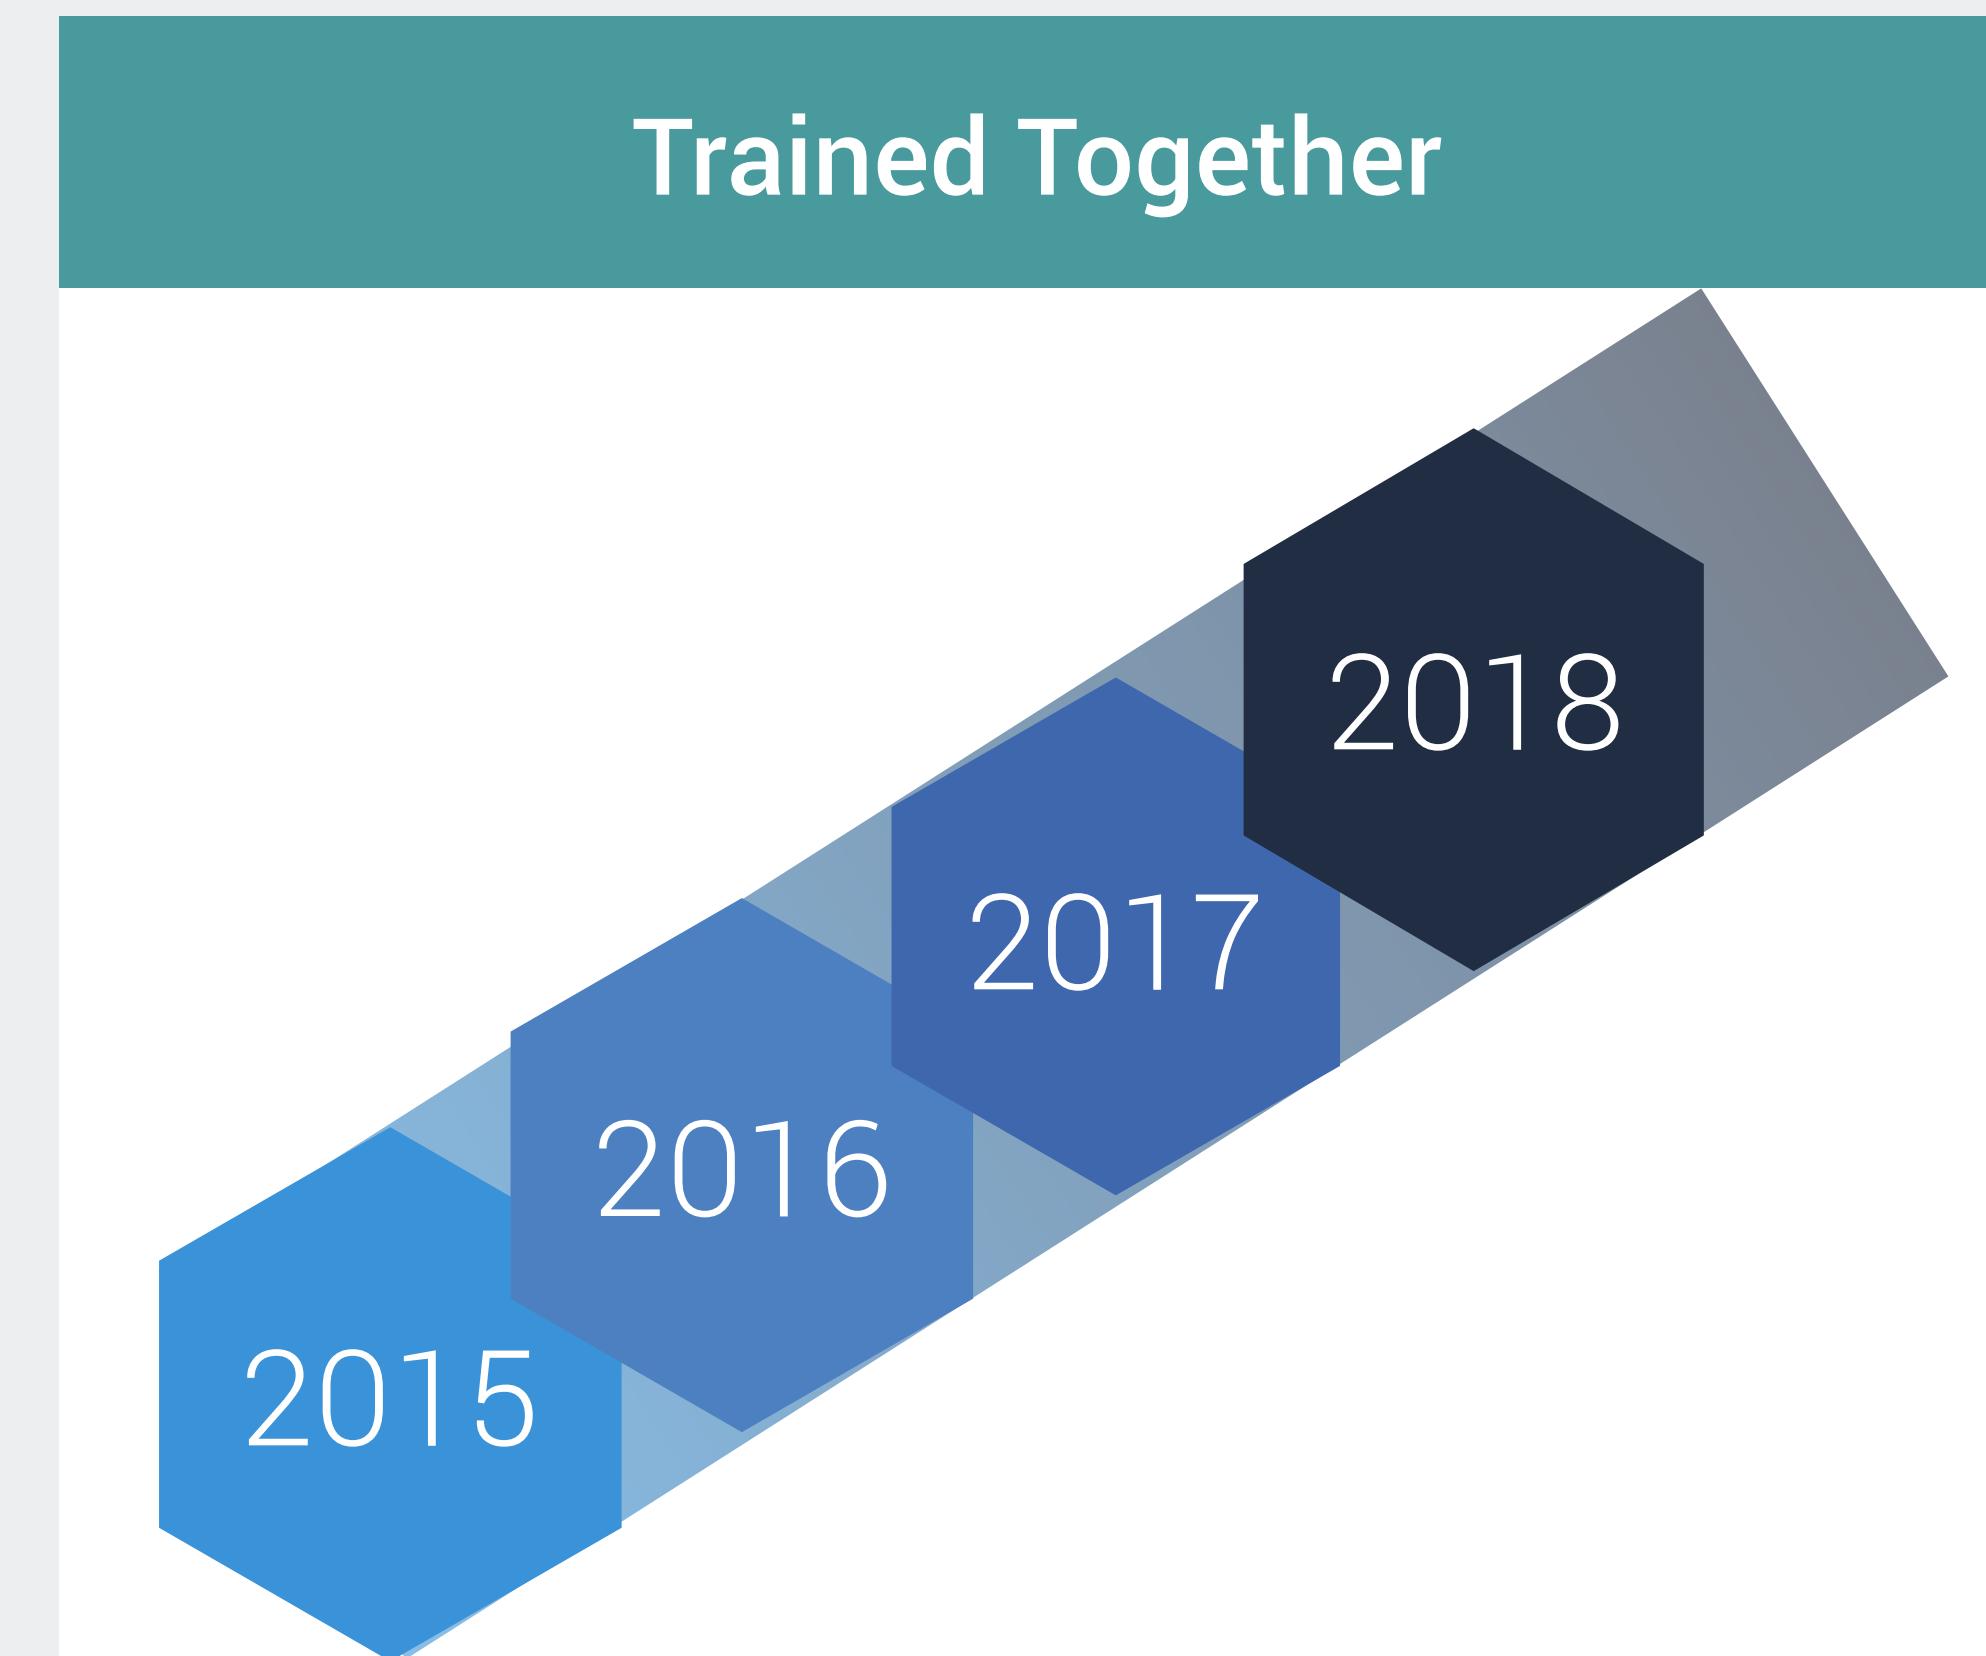
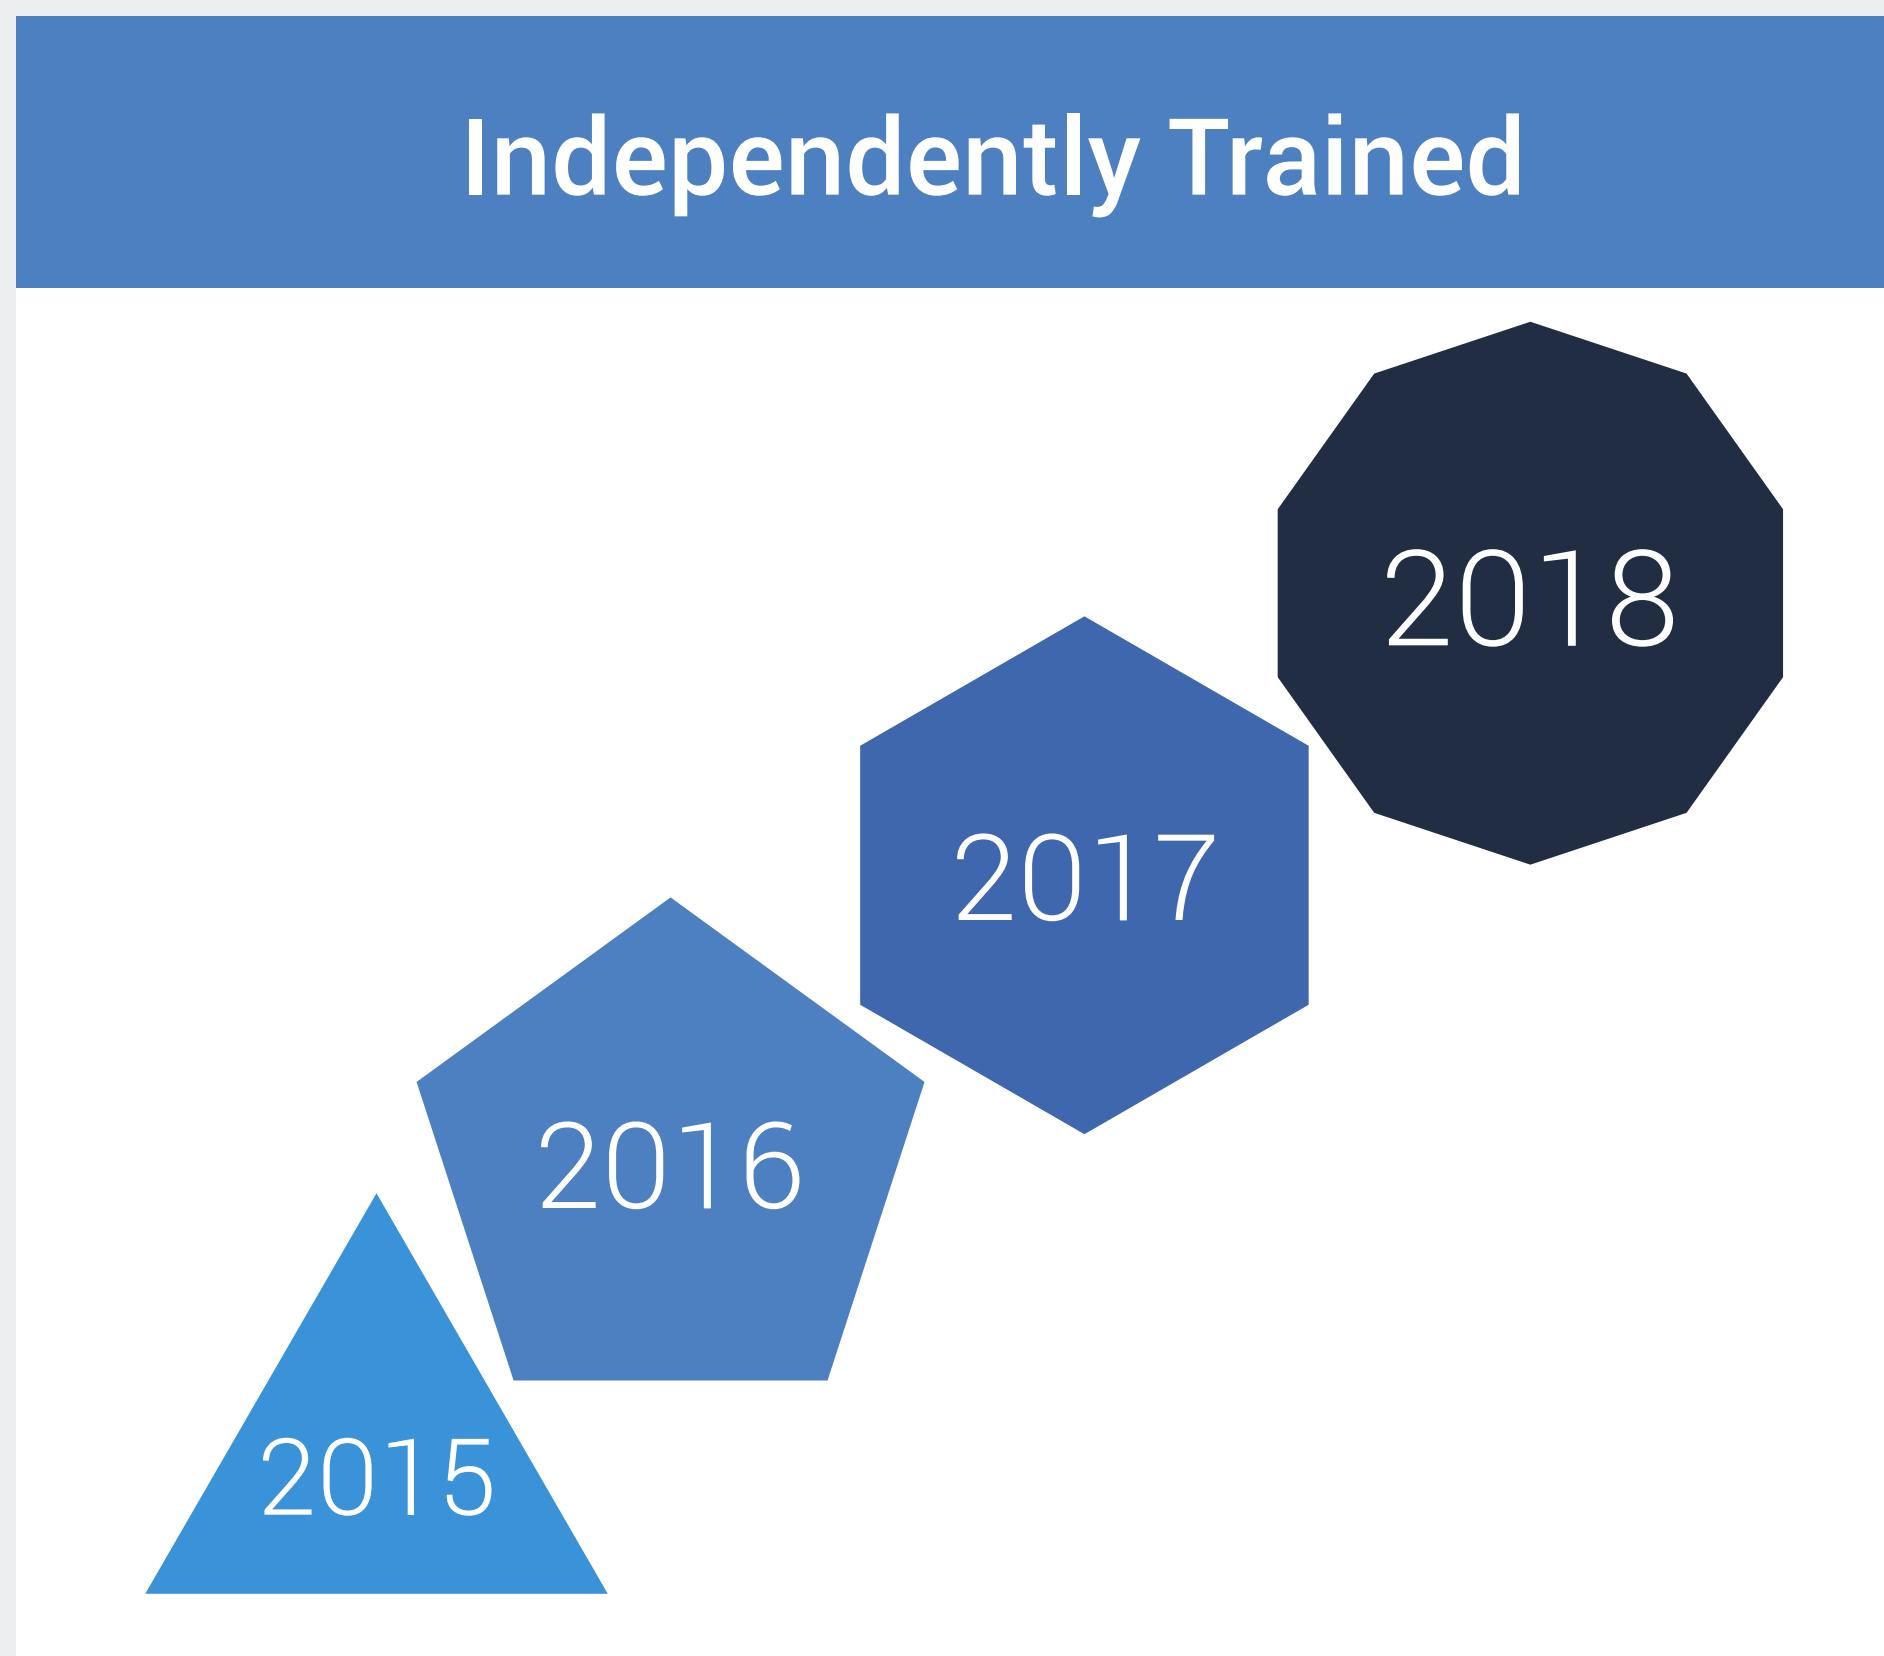


Custom embeddings identified ‘soft’ skills and language around experience



I've got 300 dimensions...
but time ain't one

Two flavors of dynamic embeddings



Kim, Chiu, Kaneki, Hedge and Petrov, [arXiv: 1405:3515](#).
Kulkarni, Al-Rfou, Perozzi and Skiena, [arXiv: 1411:3315](#).
Hamilton, Leskovec and Jurafsky, [arXiv: 1605:09096](#).

Rudolph and Blei, [arXiv: 1703.08052](#).

The benefits of dynamic Bernoulli embeddings

Dynamic embeddings

Data hungry: Sufficient data for each time slice for a quality embedding.

Requires stitching: Each time slice is trained independently, therefore dimensions are not comparable across slices.

Dynamic Bernoulli embeddings

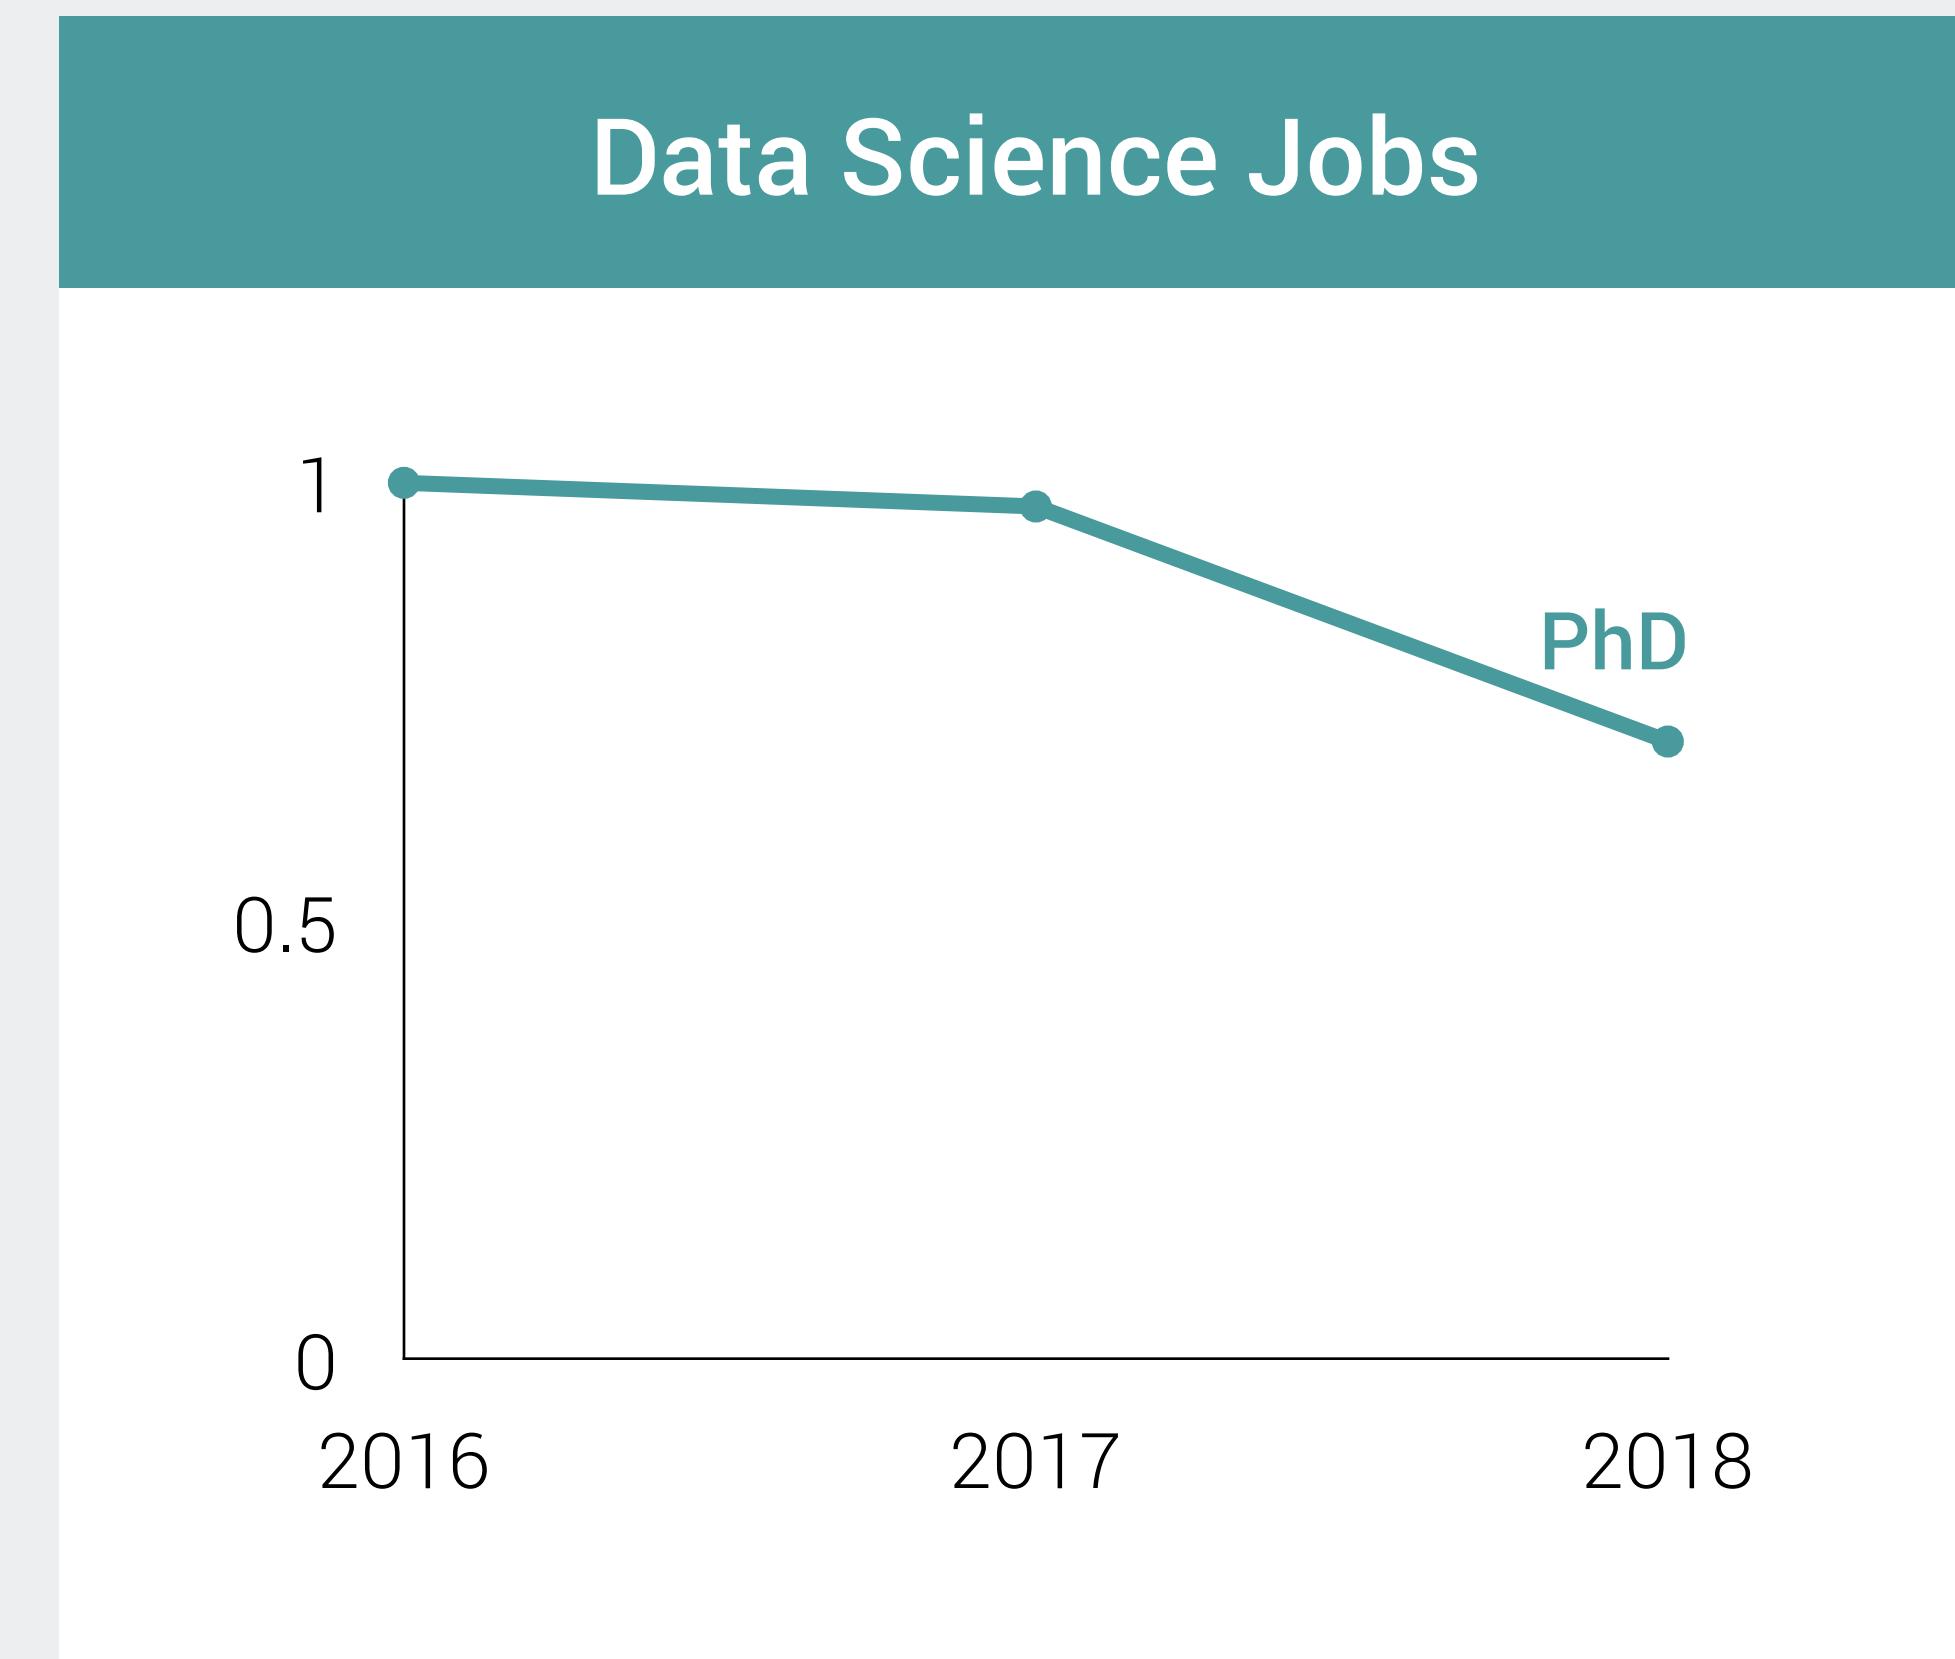
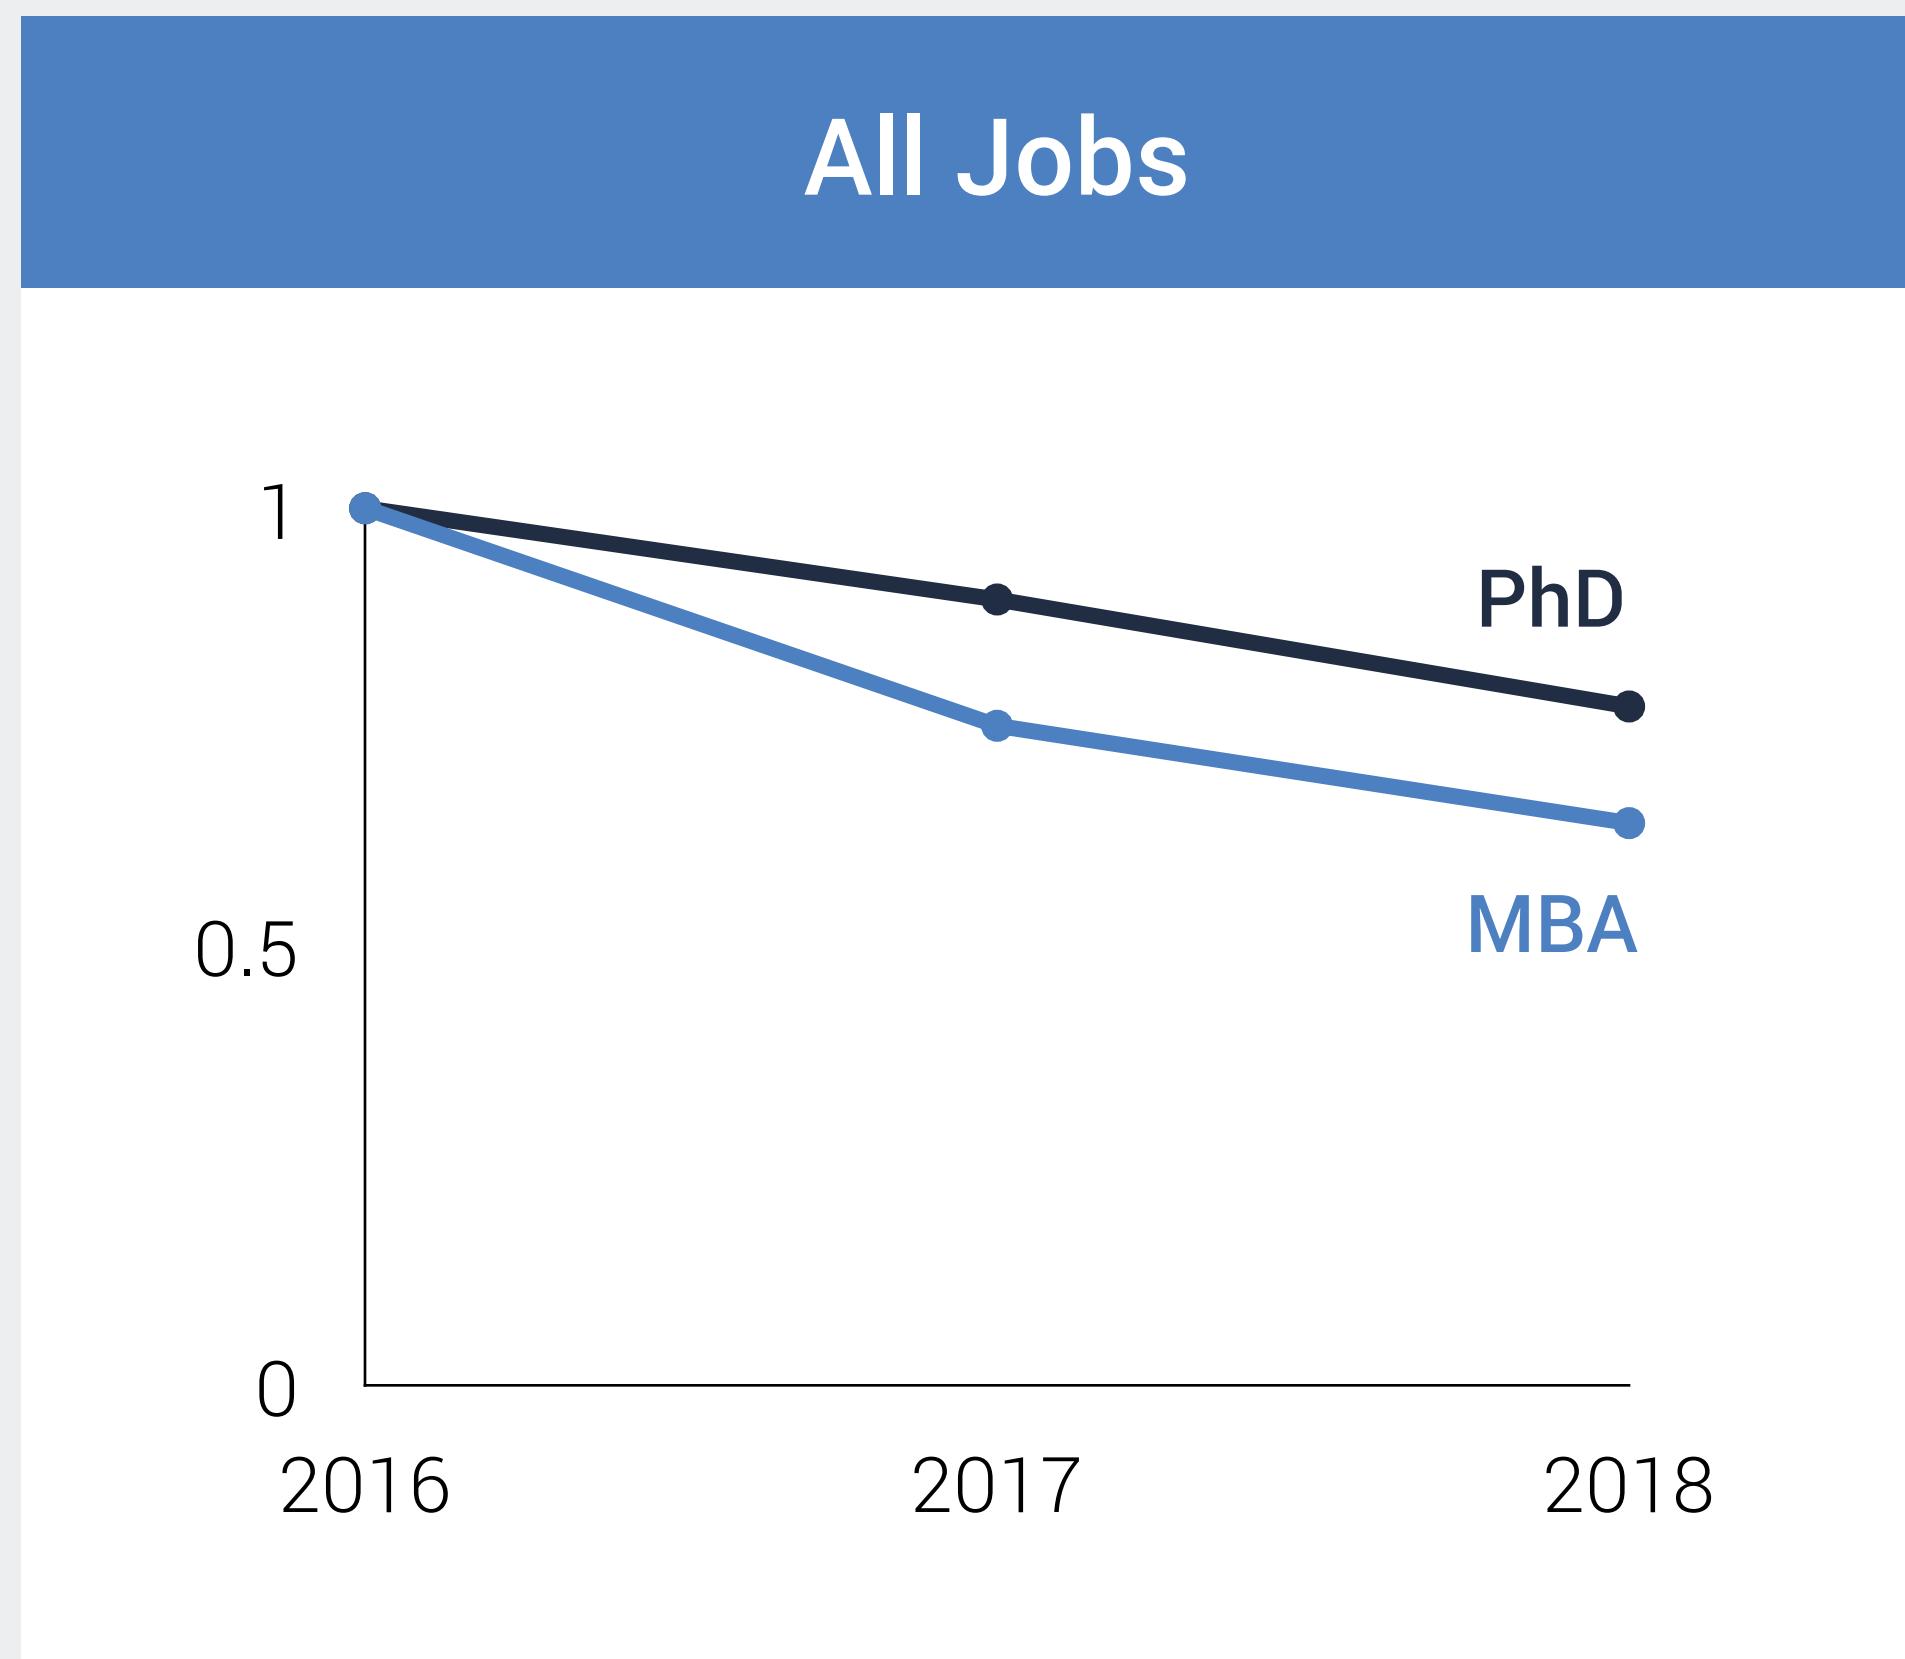
Data efficient: Treats each time slice as a sequential latent variable, enabling time slices with sparse data.

Does not require stitching/alignment: Treating time slice as a variable ensures embeddings are connected across slices.

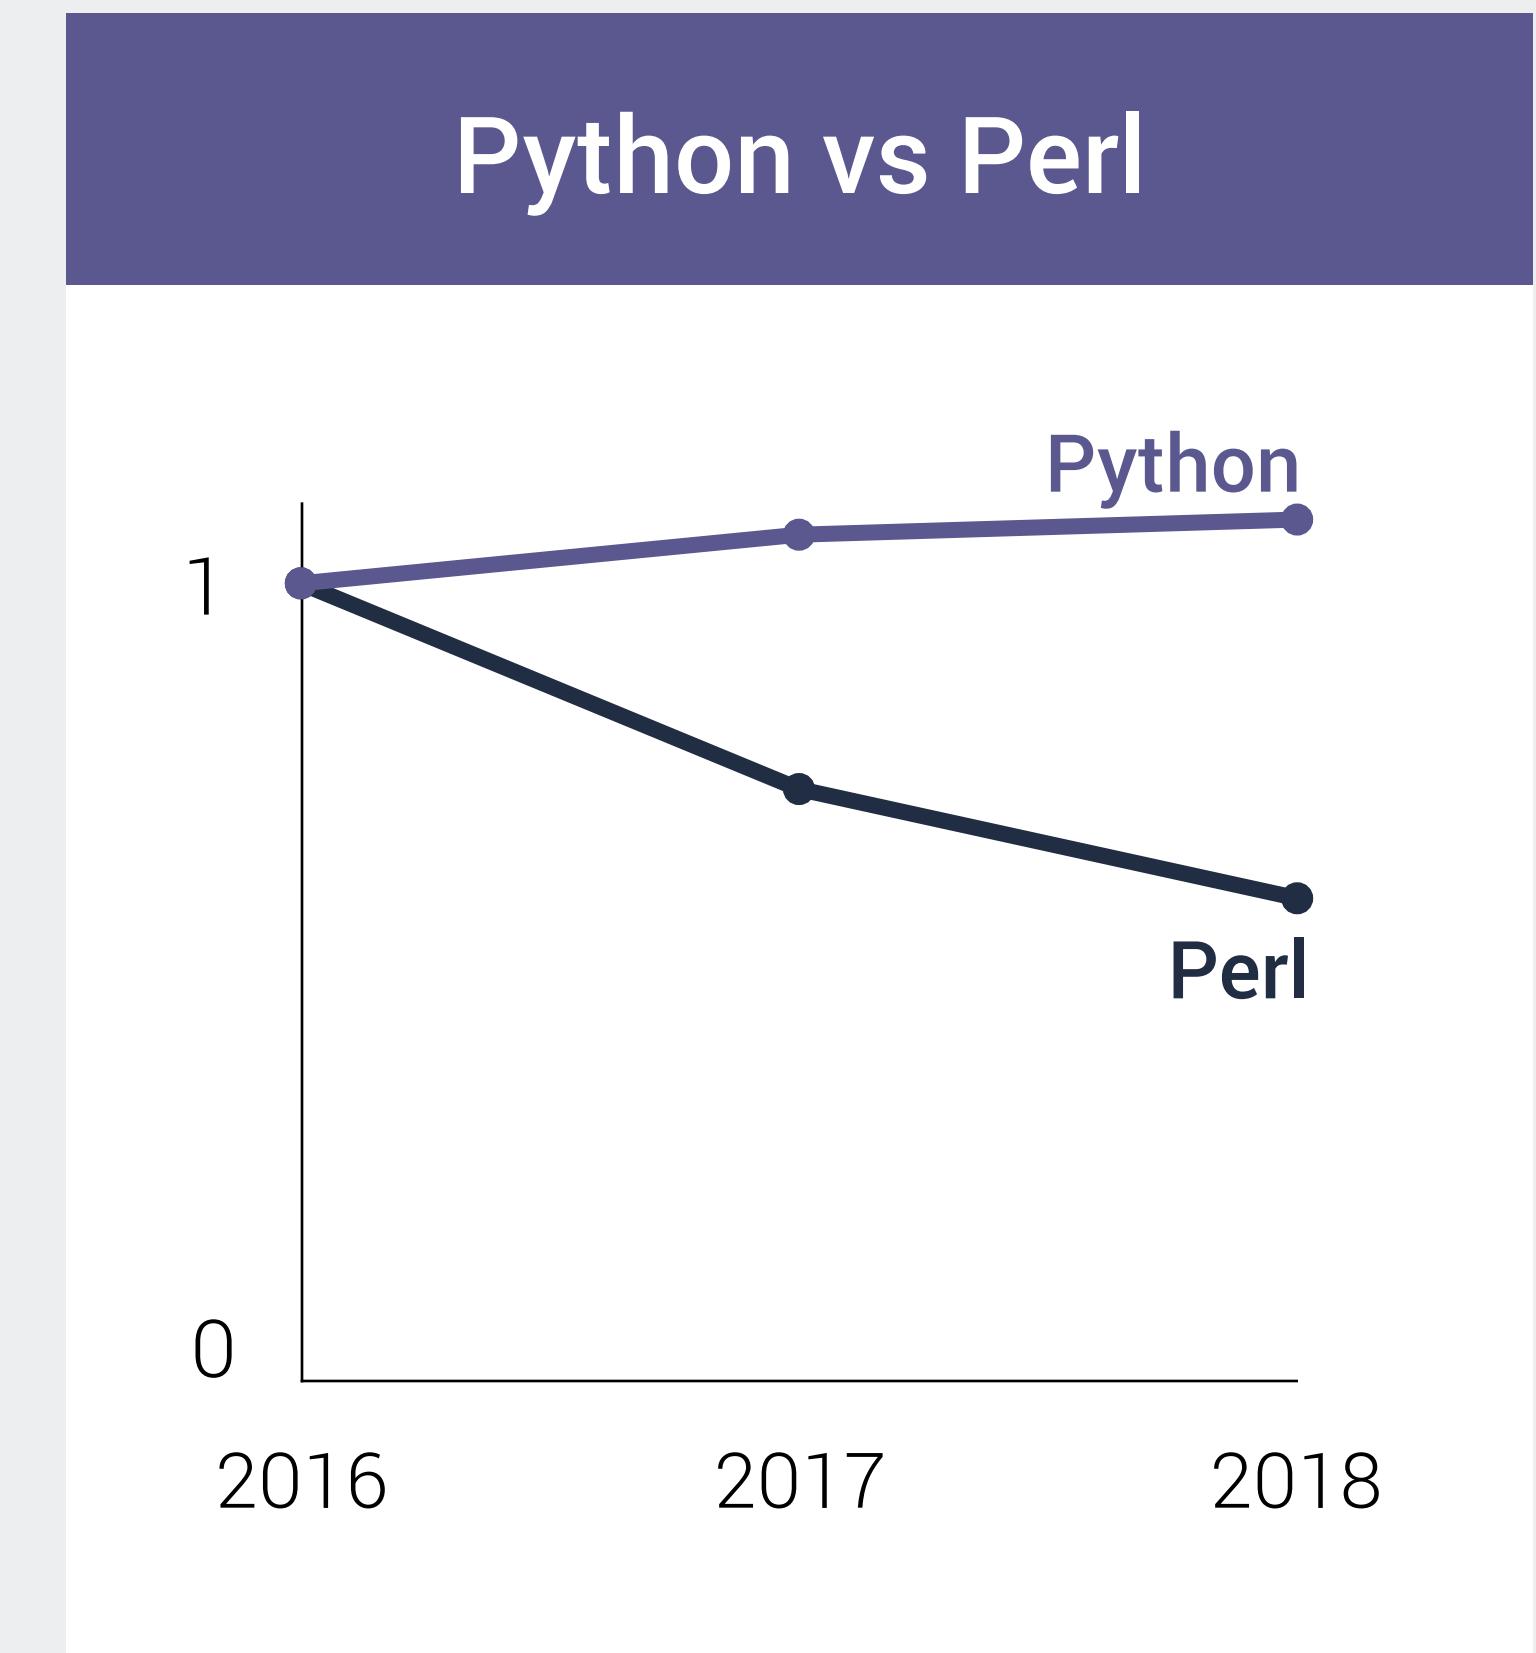
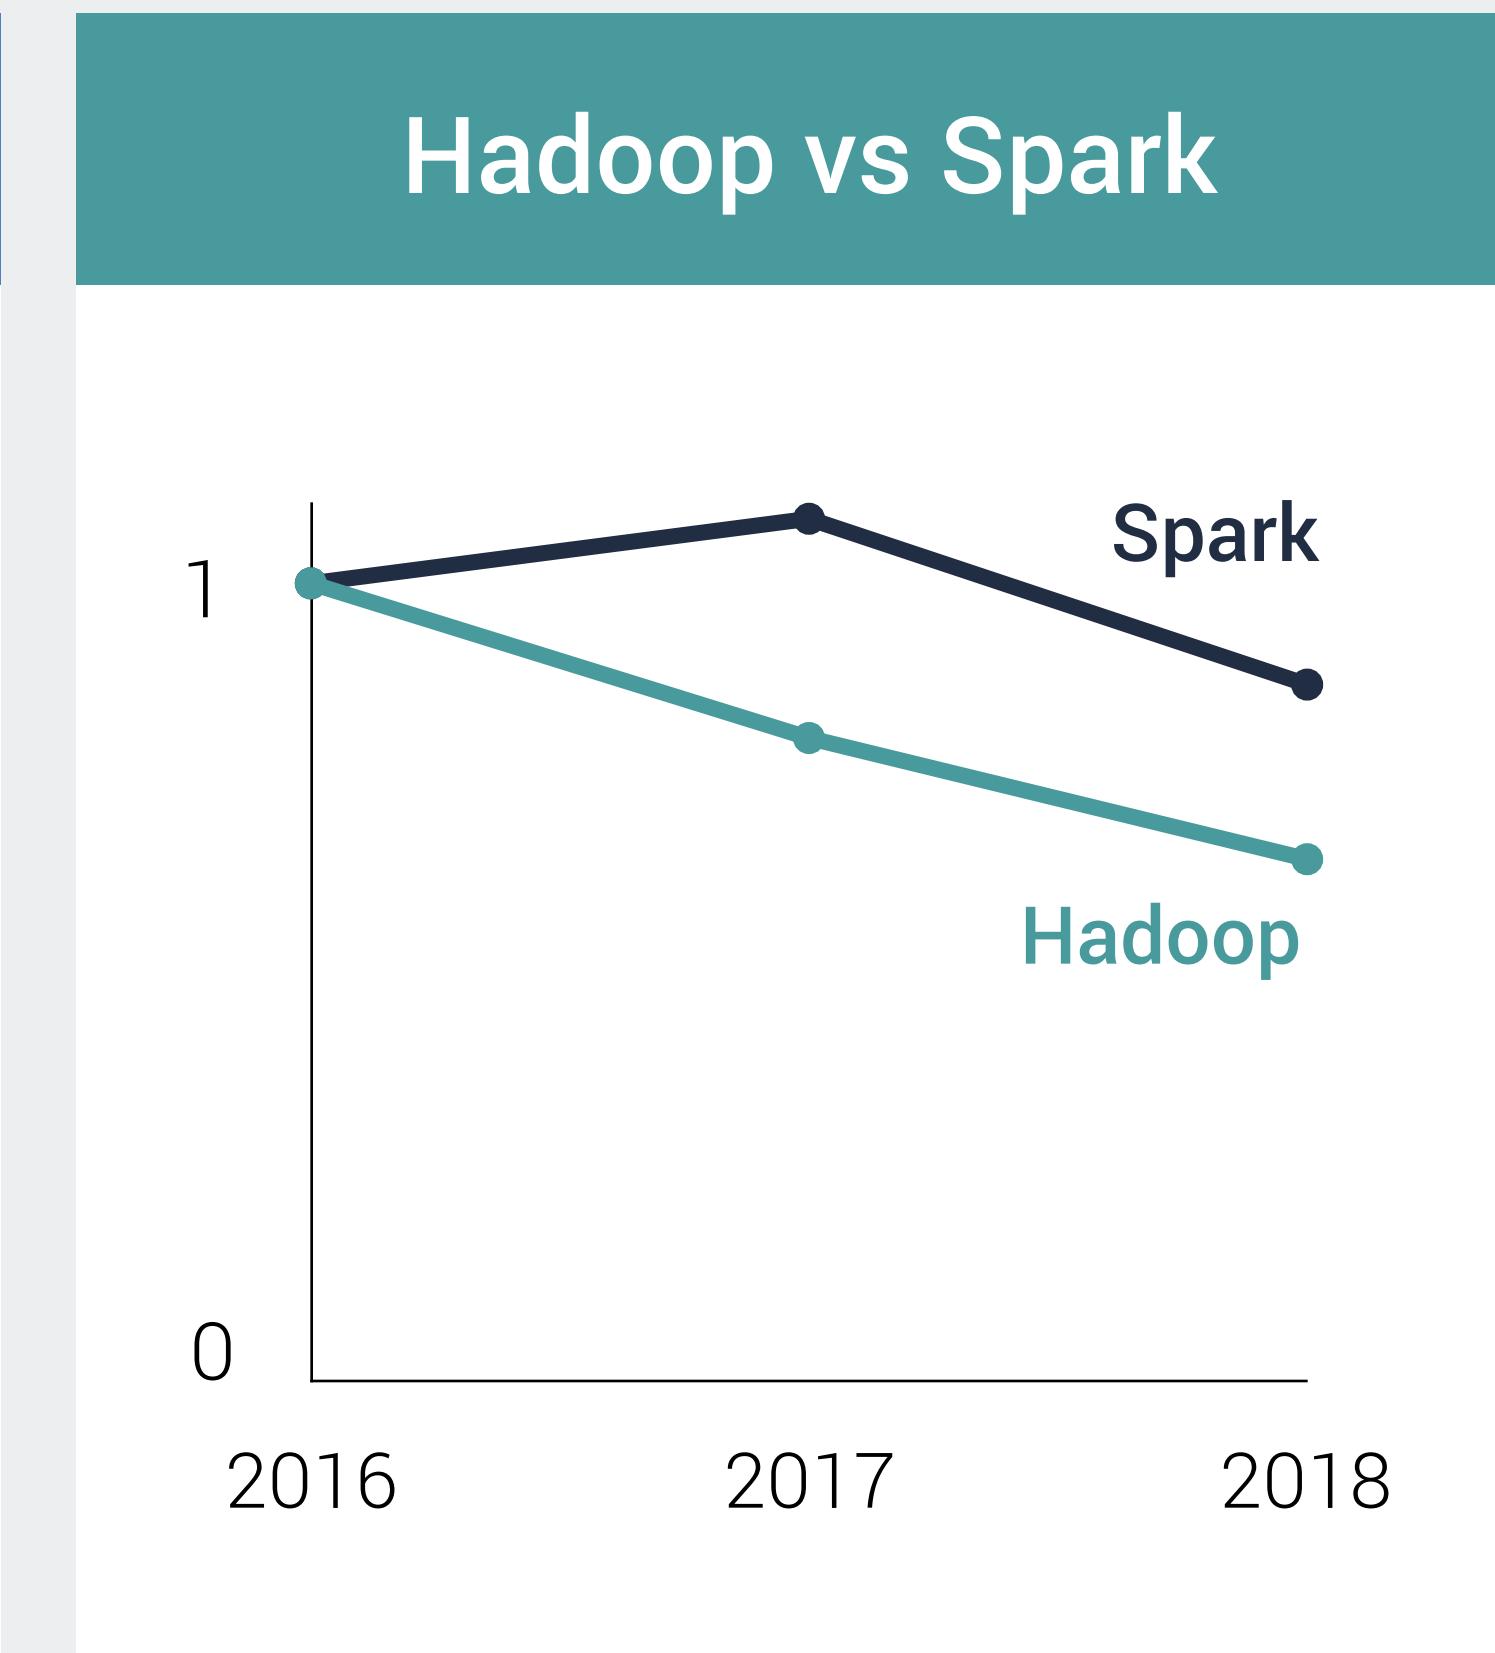
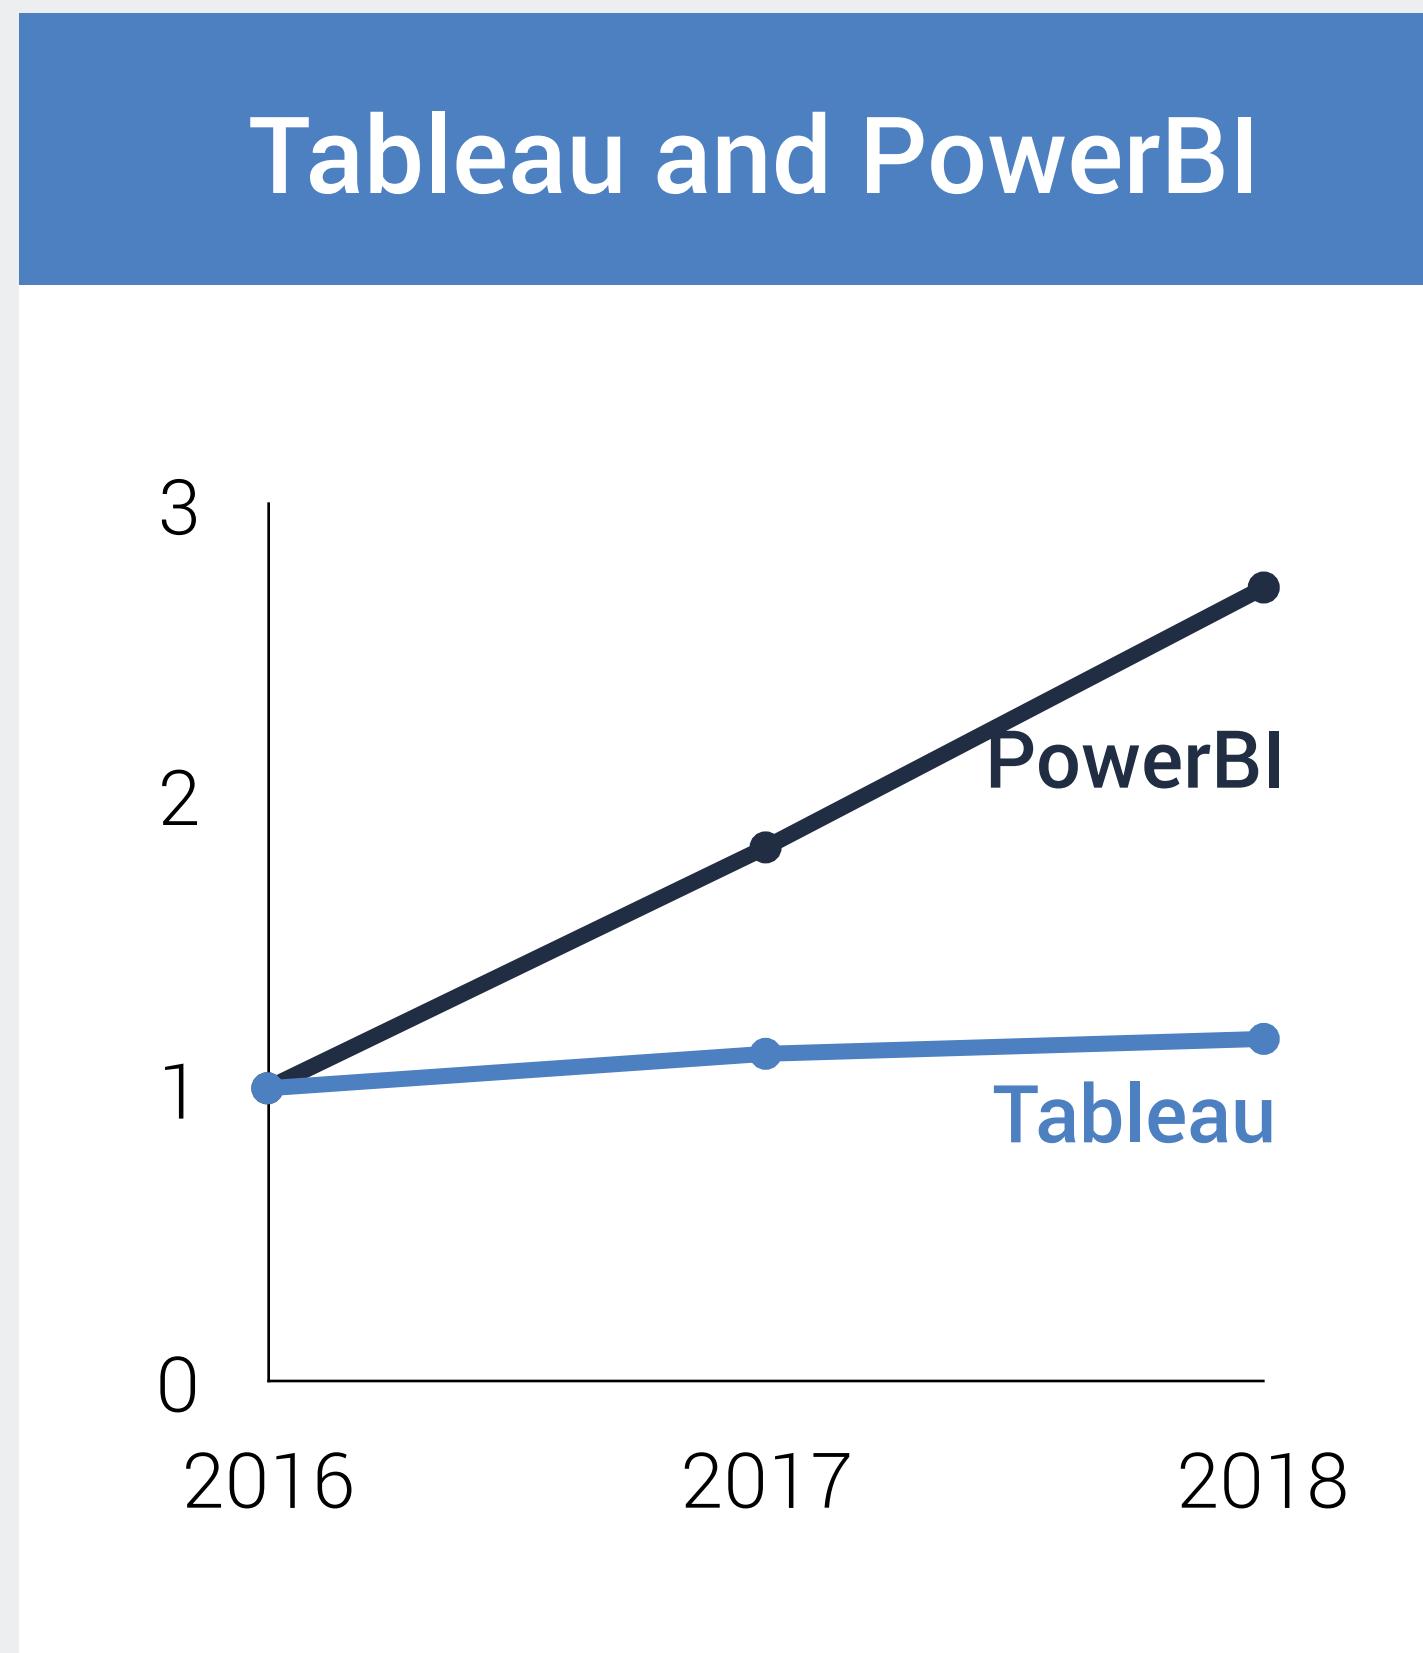
Kim, Chiu, Kaneki, Hedge and Petrov, [arXiv: 1405:3515](#).
Kulkarni, Al-Rfou, Perozzi and Skiena, [arXiv: 1411:3315](#).
Hamilton, Leskovec and Jurafsky, [arXiv: 1605:09096](#).

Balmer and Mandt, [arXiv: 1702:08359](#).
Rudolph and Blei, [arXiv: 1703:08052](#).

Demand for MBAs and PhDs falling



Data Science skills showing significant shifts



regression :: Generalized Linear Models as
word2vec :: Exponential Family Embeddings

Members of the Exponential Family of Embeddings

Bernoulli Embedding	Poisson Embedding	Gaussian Embedding
Binary Data	Count or Ordinal Data	Continuous Data
Proficiency Context programming Python	Mini Bagels Context Cream cheese Milk	JFK-CDG Context LGA-DCA JFK-DFW
Datapoint Java Context C++	Datapoint Coffee Context Orange Juice	Datapoint LAX-JFK Context LAX-LGA

Poisson embeddings capture item similarities from shopper behavior

Maruchan chicken ramen

Maruchan creamy chicken ramen
Maruchan oriental flavor ramen
Maruchan roast chicken ramen

Yoplait strawberry yogurt

Yoplait apricot mango yogurt
Yoplait strawberry orange smoothie
Yoplait strawberry banana yogurt

High Inner Product Combinations

Old Dutch potato chips & Budweiser Lager beer
Lays potato chips & DiGiorno frozen pizza

Low Inner Product Combinations

General Mills cinnamon toast & Tide Plus detergent
Beef Swanson Broth soup & Campbell Soup cans

How have data science skills changed over time?

- Flavors of static word embeddings: The Corpus Issue
- Considerations for developing custom embedding models
- Flavors of dynamic models: Dynamic Bernoulli embeddings
- Other members of the Exponential Family of Embeddings

Thank you PyData NYC!

Maryam Jahanshahi Ph.D.
Research Scientist
TapRecruit



tapRecruit.co

<http://bit.ly/pydatanyc-emb>